

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer who builds and ships agentic AI solutions for real business problems — the core of AHEAD's AI Services practice. Delivered two production multi-agent systems deployed on GCP Cloud Run with real users, a RAG concierge with controlled retrieval experiments, and an OpenEnv RL evaluation environment validated at 0.96/1.0 for Meta × Hugging Face. Three years of enterprise software engineering across BMW, healthcare, and industrial domains adds the client-facing delivery discipline that consulting AI work demands. Solid command of the full stack: LangChain, LangGraph, Google ADK, FastAPI, Docker, Cloud Run, vector DBs, and Gemini/OpenAI APIs.

CORE COMPETENCIES

Enterprise AI Solutions Delivery

Multi-Agent Orchestration (Google ADK)

RAG Pipelines & Vector DB Design

LLM API Integration (Gemini, OpenAI)

Production AI Deployment (Cloud Run)

LangChain & LangGraph

FastAPI & Docker CI-CD

Client-Facing AI System Design

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Delivered end-to-end ML systems for BMW service centres — from data validation and model design through production deployment and client-facing monitoring — achieving **80% accuracy improvement** in automated invoice generation.
- Built production backend systems across healthcare, mobility, and industrial client domains; operated in regulated, high-volume environments with strict GDPR and data governance requirements.
- Implemented structured observability and automated failure detection; reduced recurring production failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Enterprise Agentic AI Gen AI Academy APAC · 2026

Architected and shipped a multi-agent system from scratch: Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, real-time tool-use via MCP servers, parallel agent execution, AlloyDB AI persistent memory. Deployed on Cloud Run with real users. Directly models the agentic AI platform delivery AHEAD builds for enterprise clients.
Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — End-to-End Client AI Delivery

Gen AI Academy APAC · 2026

Took a conversational multi-agent AI from prototype to production — Gemini 2.5 Flash powering invoice creation, payment tracking, and analytics. Diagnosed and resolved a critical production LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Full-stack ownership: React + Vite frontend, Node.js backend, Docker, Cloud Run.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

TravelAI — RAG Pipeline with Retrieval Evaluation

Kaggle Capstone · 2025

End-to-end RAG implementation: ingestion pipelines, chunking and embedding policies, hybrid vector search, LLM grounding, and controlled retrieval experiments measuring relevance and consistency across configurations. Demonstrates the retrieval engineering and evaluation capability central to building AI knowledge systems for enterprise clients.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Python

WorkflowEnv — LLM Agent Evaluation Framework

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents — FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, strict telemetry logging. Both phases validated. **Avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · LLM agent evaluation · telemetry

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

LLM & Agents: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, multi-agent orchestration, ReAct pattern, prompt engineering, LLM evaluation

AI Systems: RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, RL environment design, OpenEnv

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, MLflow, Datadog

Languages: Python, FastAPI, Node.js, TypeScript, SQL, Bash, C#