

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer specialising in RAG pipelines, LLM evaluation, and production-grade conversational AI — the core of Amira Learning's Gen AI stack. Built an OpenEnv-compliant LLM evaluation framework with 7-criterion partial-credit grading and hallucination-aware scoring, validated at 0.96/1.0 (Meta × Hugging Face 2026). Built an end-to-end RAG system with controlled retrieval experiments across chunking, embedding, and retrieval configurations. Shipped two production multi-agent AI systems on GCP Cloud Run with real users. Strong command of LangChain, LangGraph, RAG pipelines, vector DBs, and prompt engineering for LLM reliability — the exact skills required to build conversational AI that children trust and that educators can rely on.

CORE COMPETENCIES

RAG Pipelines & Retrieval Evaluation

LLM Evaluation & Hallucination Mitigation

Production Conversational AI Systems

LangChain & LangGraph

Vector DB Tuning (Pinecone)

Prompt Engineering & Guardrails

Multi-Agent Orchestration (Google ADK)

Docker & GCP Cloud Run

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven data validation and automation systems for BMW service centres — end-to-end pipeline ownership from feature engineering through deployment and monitoring, delivering **80% accuracy improvement**.
- Implemented structured observability and automated failure detection across healthcare, mobility, and industrial domains — the production reliability mindset that safety-critical AI for children demands.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — LLM Evaluation & Guardrail Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment for LLM agents — 7-criterion partial-credit grader covering task completion, instruction-following, constraint adherence, output grounding, and hallucination signals. Full telemetry logging, automated scoring pipeline, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0, 3/3 tasks passing**. Directly demonstrates the LLM evaluation, guardrail engineering, and hallucination mitigation skills Amira's Gen AI stack requires.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · guardrails · telemetry · automated scoring

TravelAI — RAG Pipeline with Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG system with heterogeneous data ingestion, chunking and embedding pipelines, hybrid vector search, LLM grounding, and controlled experiments measuring retrieval relevance and response consistency across embedding models and retrieval configurations. Demonstrates the RAG architecture depth and retrieval evaluation rigour Amira's conversational AI requires.

RAG · Vector DB · Embeddings · Pinecone · Semantic search · LLM grounding · Retrieval evaluation · Python

AI Invoice Agent — Production Conversational LLM Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI from prototype to production — Gemini 2.5 Flash powering intent routing, invoice creation, payment tracking, and analytics. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** FastAPI + Node.js backend, Docker, Cloud Run.

Gemini 2.5 Flash · LangChain · FastAPI · Node.js · Docker · Cloud Run

Multi-Agent Productivity Assistant — Production Agentic System Gen AI Academy APAC · 2026

Architected multi-agent AI from scratch: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel execution, real-time MCP tool-use, AlloyDB AI persistent memory across sessions. Deployed on Cloud Run with real users.

Gemini 2.0 Flash · Google ADK · LangGraph · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

RAG & Eval: RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM evaluation, hallucination mitigation, guardrails, automated scoring, retrieval evaluation, LlamaIndex

LLMs & Agents: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, prompt engineering, agent orchestration, conversational AI, multi-turn dialogue

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, MLflow, Datadog

Languages: Python, FastAPI, Node.js, SQL, Bash