Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer specialising in retrieval systems, LLM orchestration, prompting, and evaluation — the full applied AI stack Auctor is building for professional services. Built an end-to-end RAG system with controlled retrieval experiments: heterogeneous data ingestion, chunking, hybrid vector search, LLM grounding, and systematic evaluation across embedding models and retrieval configurations (TravelAI, Kaggle 2025). Shipped an RL evaluation environment with 7-criterion partial-credit grader and automated scoring validated at 0.96/1.0 (WorkflowEnv, Meta × Hugging Face 2026). Production multi-agent orchestration with MCP tool-use deployed on Cloud Run with real users. Three years of enterprise backend engineering adds the reliability foundation professional services AI demands.

CORE COMPETENCIES

- RAG Pipelines & Document Retrieval
- LLM Orchestration & Tool Calling
- Prompt Engineering & Context Management
- LLM Evaluation Framework Design
- Multi-Agent Orchestration (Google ADK)
- LangChain & LangGraph
- FastAPI & Python Backend
- Production LLM Reliability

WORK EXPERIENCE

PICA GmbH
2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML backend systems for BMW service centres — from data ingestion pipelines and model design through production deployment and monitoring, delivering **80% accuracy improvement**.
- Built production backend systems across healthcare, mobility, and industrial domains; implemented structured observability and anomaly detection — the reliability engineering discipline professional services AI requires.

AAM IT GmbH / Vorwerk Switzerland
Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes
May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

TravelAI — RAG Pipeline & Document Retrieval System
Kaggle Capstone · 2025

End-to-end RAG system directly applicable to Auctor's document understanding and retrieval work: heterogeneous data ingestion, chunking and embedding pipeline design, hybrid vector search, LLM grounding, and controlled experiments measuring retrieval relevance and response consistency across embedding models and retrieval configurations. Demonstrates the retrieval architecture depth and systematic evaluation that professional services AI demands for reliable factual output.

RAG · Vector DB · Pinecone · Embeddings · Hybrid search · LLM grounding · Retrieval evaluation · Python · FastAPI

WorkflowEnv — LLM Evaluation & Scoring Framework
OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant evaluation environment — 7-criterion partial-credit grader covering task completion, instruction adherence, constraint satisfaction, output quality, and scoring confidence. Telemetry logging for full experiment traceability. Automated scoring pipeline. Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.** Demonstrates the evaluation and evals engineering Auctor's applied AI work requires to close the feedback loop between retrieval quality and model output.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · partial-credit grading · telemetry · automated scoring

Multi-Agent Productivity Assistant — LLM Orchestration & Tool Calling in Production
Gen AI Academy APAC · 2026

Architected production multi-agent system with real-time tool-use: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel execution, MCP tool-use, AlloyDB AI persistent memory across sessions. Deployed on Cloud Run with real users. Demonstrates the LLM orchestration, multi-step tool calling, and context management skills central to Auctor's platform.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Production Conversational AI
Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI from prototype to production — Gemini 2.5 Flash powering intent routing and tool calling via LangChain. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** FastAPI + Node.js backend, Docker, Cloud Run.

Gemini 2.5 Flash · LangChain · FastAPI · Node.js · Docker · Cloud Run

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany — 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University — 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) — 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive — 2024–2025

SKILLS

Retrieval & RAG: RAG pipelines, vector DBs (Pinecone), embeddings, hybrid search, chunking strategies, LLM grounding, retrieval evaluation, document understanding, context management

LLMs & Orchestration: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, multi-agent orchestration, tool calling, prompt engineering, ReAct patterns

Eval & Observability: LLM evaluation frameworks, partial-credit grading, automated scoring, telemetry logging, MLflow, Datadog, structured observability

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, AlloyDB AI