

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer specialising in turning LLM and agent research into production-grade systems — the core of what Axiomatic AI's Applied AI Engineer role requires. Built an OpenEnv-compliant RL evaluation environment for LLM workflow agents (validated 0.96/1.0, Meta × Hugging Face 2026), two production multi-agent systems deployed on GCP Cloud Run with real users, and a multimodal AI observability copilot integrating Datadog with Vertex AI. Strong command of the full stack from prototype to deployment: LangChain, LangGraph, Google ADK, FastAPI, Docker, Cloud Run. Three years of enterprise software engineering in regulated environments adds the production discipline that bridges R&D to real-world reliability.

CORE COMPETENCIES

Research-to-Production AI Engineering

LLM Agent System Design

Multi-Agent Orchestration (Google ADK)

LLM Evaluation Framework Design

Gemini / OpenAI API Integration

GCP Cloud Run & Docker CI-CD

RAG Pipelines & Semantic Retrieval

FastAPI & Production Deployment

WORK EXPERIENCE

PICA GmbH

2018 – 2021

Software Developer · Munich, Germany

- Translated business requirements into production ML systems for BMW service centres — designing the full pipeline from data validation and model architecture through deployment and monitoring, delivering **80% accuracy improvement**.
- Built backend systems across healthcare, mobility, and industrial domains in regulated, high-volume environments — the production discipline that turns research prototypes into reliable systems.
- Implemented structured observability and automated failure detection; reduced recurring production failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland

Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes

May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

WorkflowEnv — OpenEnv RL Environment for LLM Agents

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents — bridging agent research to production evaluation infrastructure. FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, strict telemetry logging. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.** Directly demonstrates the R&D-to-production capability Axiomatic AI requires.

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry logging

Multi-Agent Productivity Assistant — Production Agentic System

Gen AI Academy APAC · 2026

Architected and shipped a multi-agent system from scratch: Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, real-time tool-use via MCP servers, parallel agent execution, AlloyDB AI persistent memory. Cloud Run deployment with real users.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — LLM Prototype to Reliable Production System

Gen AI Academy APAC · 2026

Took a conversational multi-agent AI from prototype to production — Gemini 2.5 Flash powering invoice creation, payment tracking, and analytics. Diagnosed and resolved a critical reliability issue (API version mismatch) in a live system. **Impact: ~70% faster workflows.** Full-stack ownership.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

AI SRE Copilot — Multimodal LLM Application

DEVPOST Hackathon · 2025

Multimodal AI application integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis. FastAPI on Cloud Run. Demonstrates LLM workflow design for complex operational data inputs.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

- LLM & Agents: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, multi-agent orchestration, ReAct pattern, prompt engineering, LLM evaluation
- AI Systems: RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, RL environment design, OpenEnv
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, MLflow, Datadog
- Languages: Python, FastAPI, Node.js, TypeScript, SQL, Bash