

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer with a proven record of taking LLM-powered systems from zero to production — two agentic AI systems shipped to real users on GCP Cloud Run in 2026, an OpenEnv RL evaluation environment validated at 0.96/1.0 for Meta × Hugging Face, and a production LLM reliability issue resolved mid-flight in a live system. Deep command of LLM API integration (Gemini, OpenAI), multi-agent orchestration, and end-to-end AI product delivery. The kind of engineer who builds the thing, ships it, and fixes it when it breaks in production. Available for worldwide async-first remote collaboration.

CORE COMPETENCIES

Zero-to-One LLM Product Delivery

LLM API Integration (Gemini, OpenAI)

Multi-Agent Orchestration (Google ADK)

Agentic AI System Design

Production LLM Reliability Engineering

LLM Evaluation Framework Design

GCP Cloud Run & Docker CI-CD

FastAPI & Full-Stack AI Deployment

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML systems for BMW service centres — from data pipelines and model design through production deployment and monitoring — delivering **80% accuracy improvement** in automated invoice generation.
- Built production backend systems across healthcare, mobility, and industrial domains; operated in regulated environments with strict data governance requirements.
- Strengthened production observability via structured error handling and logging; reduced recurring failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk’s Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Zero to Production Gen AI Academy APAC · 2026

Built from zero to production: Orchestrator Agent routes intent to four specialist sub-agents (Calendar, Task, Notes, Email) using Gemini 2.0 Flash + Google ADK. Real-time tool-use via Google Calendar MCP and Gmail MCP. Parallel agent execution resolves multi-step user requests in a single turn. AlloyDB AI persistent memory. Deployed on Cloud Run with real users.
Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Shipped & Reliability-Tested in Production

Gen AI Academy APAC · 2026

Designed and shipped a conversational AI agent from scratch — Gemini 2.5 Flash powering invoice creation, payment tracking, and real-time analytics. Resolved a critical LLM reliability issue in production (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Full-stack ownership: React + Vite frontend, Node.js backend, Docker, Cloud Run.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

WorkflowEnv — OpenEnv RL Environment for LLM Agents

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents — FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, strict telemetry logging. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation

AI SRE Copilot — Multimodal LLM Application

DEVPOST Hackathon · 2025

Multimodal AI application integrating Datadog with Vertex AI (Gemini) for root-cause analysis. FastAPI + Cloud Run deployment. Demonstrates LLM workflow design for complex, data-rich inputs.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany

2017 – 2021
- Thesis: IEEE-published — large-scale ML pipeline for geospatial data processing adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University

2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)

2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive

2024–2025

SKILLS

- LLM APIs: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, prompt engineering, LLM evaluation, fine-tuning
- AI Systems: Multi-agent orchestration, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, ReAct pattern
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, MLflow, Datadog
- Languages: Python, FastAPI, Node.js, TypeScript, React, SQL, Bash