

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Platform Engineer with hands-on experience building production LLM orchestration systems with MCP integrations, agentic SDKs, and multi-agent backends — the exact foundation CaptivateIQ's Agentic SDK team is building. Shipped a production multi-agent system with real-time MCP tool-use (Google Calendar MCP, Gmail MCP), parallel agent execution via Google ADK + LangGraph, and persistent session memory — deployed on Cloud Run with real users. Built a FastAPI-based LLM evaluation environment validated at 0.96/1.0, demonstrating the API-first design and developer platform thinking the Agentic SDK role requires. Three years of enterprise backend engineering adds the systems discipline that internal developer tooling demands.

CORE COMPETENCIES

LLM Orchestration & Agentic SDK Design

MCP Tool Integration (Real-Time)

LangChain & LangGraph

Multi-Agent Architecture (Google ADK)

FastAPI & Python Backend Systems

Vector DB & Knowledge Systems

React / TypeScript Frontend

Docker & GCP Cloud Run CI-CD

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML backend systems for BMW service centres — from data pipelines and model integration through production deployment and monitoring — delivering **80% accuracy improvement**.
- Built production backend systems across healthcare, mobility, and industrial domains; implemented structured observability and automated failure detection to maintain reliability under high-volume enterprise workloads.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — MCP + LangGraph Agentic System Gen AI Academy APAC · 2026

Architected and shipped a production agentic system with real-time MCP tool integrations — the exact architecture CaptivateIQ's Agentic SDK enables for product teams. Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, parallel execution resolving multi-step requests in a single turn, Google Calendar MCP + Gmail MCP live tool-use, AlloyDB AI persistent session memory. Deployed on Cloud Run with CI/CD. Real users in production.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

WorkflowEnv — FastAPI LLM Evaluation Platform

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed a FastAPI-based RL evaluation platform for LLM agents — 7-criterion partial-credit grader, telemetry logging, Docker deployment to HF Spaces. API-first design used by external agents to submit results. Both phases validated. **Avg score 0.96/1.0.** Demonstrates developer platform API design and internal SDK thinking directly applicable to CaptivateIQ's Agentic SDK.

FastAPI · Docker · HF Spaces · OpenEnv · API design · LLM evaluation · telemetry

AI Invoice Agent — Full-Stack LLM Application

Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI from prototype to production — React + Vite frontend, FastAPI + Node.js backend, Gemini 2.5 Flash LLM orchestration via LangChain. Full-stack delivery across the same stack CaptivateIQ uses. Diagnosed and resolved a critical LLM reliability issue mid-flight. **Impact: ~70% faster workflows.**

Gemini 2.5 Flash · LangChain · React + Vite · FastAPI · Node.js · Docker · Cloud Run

TravelAI — RAG & Knowledge System

Kaggle Capstone · 2025

End-to-end RAG system: ingestion pipelines, chunking and embedding policies, vector search, LLM grounding, and controlled retrieval experiments. Demonstrates the knowledge system and vector DB integration design relevant to CaptivateIQ's Agentic SDK knowledge layer.

RAG · Vector DB · Pinecone · Embeddings · Semantic search · LLM grounding · Python

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

AI Platform: LLM orchestration, LangChain, LangGraph, Google ADK, MCP tool integration, agentic SDK design, multi-agent architecture, prompt engineering, vector DBs (Pinecone)

Full-Stack: Python, FastAPI, React, TypeScript, Node.js, SQL, REST APIs, GraphQL familiarity

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, AlloyDB AI, HF Spaces, Datadog

LLM APIs: Gemini 2.5/2.0, OpenAI API, Anthropic API, RAG pipelines, embeddings, semantic search