

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer specialising in agentic workflow design and production LLM systems — the core of Cohere's Applied AI Engineer role. Shipped two agentic AI systems from prototype to production on GCP Cloud Run: a multi-agent orchestrator with parallel sub-agent execution, real-time MCP tool-use, and persistent session memory; and a conversational multi-agent AI delivering **~70% workflow acceleration** in production. Strong command of the full agentic stack — ReAct patterns, multi-step tool calling, RAG grounding, LLM evaluation, and performance optimisation — built with LangGraph, LangChain, Google ADK, and Gemini/OpenAI APIs. Three years of enterprise backend engineering in high-stakes regulated domains adds the reliability discipline that enterprise-grade agentic workflows demand.

CORE COMPETENCIES

Agentic Workflow Design & Delivery

Multi-Step Tool Calling & ReAct Patterns

LangChain & LangGraph

RAG Pipelines & LLM Grounding

LLM Evaluation & Performance Optimisation

Multi-Agent Orchestration (Google ADK)

Production LLM Reliability

FastAPI & GCP Cloud Run

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML backend systems for BMW service centres — from data pipelines and model design through production deployment and monitoring — delivering **80% accuracy improvement**.
- Built production systems across healthcare, mobility, and industrial domains in regulated environments; implemented structured observability to maintain reliability at scale — the engineering rigour that enterprise AI workflows require.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Production Agentic Workflow System Gen AI Academy APAC · 2026

Designed and shipped a production agentic workflow system from scratch: Orchestrator Agent routing intent and decomposing multi-step requests across four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK. Parallel agent execution resolves complex requests in a single turn. Real-time tool-use via MCP servers. AlloyDB AI maintains context across sessions. Cloud Build CI/CD, Cloud Run deployment. Real users in production — demonstrates the full agentic workflow lifecycle Cohere's Applied AI Engineers own.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Conversational Agentic AI in Production

Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI with intent routing, tool calling, and analytics — Gemini 2.5 Flash + LangChain on FastAPI + Node.js. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) in a live system mid-flight. **Impact: ~70% faster workflows.** Full-stack: React + Vite frontend, Docker, Cloud Run. End-to-end ownership from prototype to production.

Gemini 2.5 Flash · LangChain · React + Vite · FastAPI · Node.js · Docker · Cloud Run

TravelAI — RAG Pipeline with Retrieval Evaluation

Kaggle Capstone · 2025

End-to-end RAG system: ingestion pipelines, chunking and embedding, hybrid vector search, LLM grounding, and controlled experiments measuring retrieval relevance and response quality across configurations. Demonstrates the retrieval evaluation and RAG grounding discipline that enterprise-grade agentic workflows require for reliable factual output.

RAG · Vector DB · Pinecone · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

WorkflowEnv — LLM Agent Evaluation Framework

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL evaluation environment for LLM workflow agents — 7-criterion partial-credit grader, telemetry logging, Docker deployment. Both phases validated. **Avg score 0.96/1.0.** Demonstrates systematic evaluation and performance optimisation — the feedback loop that makes agentic systems reliable.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · telemetry · RL environment design

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

- Agentic AI: LangChain, LangGraph, Google ADK, multi-agent orchestration, ReAct pattern, MCP tool-use, parallel execution, multi-step reasoning, tool calling, failure handling
- LLMs & RAG: Gemini 2.5/2.0, OpenAI API, Anthropic API, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM evaluation, prompt engineering
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, AlloyDB AI, HF Spaces, Datadog
- Languages: Python, FastAPI, Node.js, TypeScript, React, SQL, Bash