

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

Applied AI Engineer specializing in production-grade agentic workflows and LLM-powered systems. Shipped a multi-agent orchestration system with parallel execution, tool use, and persistent memory serving real users; an LLM evaluation environment scoring 0.96/1.0 with automated grading and telemetry; and an end-to-end RAG pipeline with systematic retrieval experiments. Strong command of LangGraph, Google ADK, FastAPI, and GCP -- translating ambiguous enterprise problems into reliable, observable, and auditable AI systems.

CORE COMPETENCIES

Agentic Workflow Design

LLM-Powered Agent Systems

Multi-Agent Orchestration

RAG Architectures

LLM Evaluation & Scoring

Prompt Engineering & Guardrails

Enterprise AI Delivery

LLMOps & Observability

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer | Munich, Germany

- Built ML-driven data validation and automation pipelines for BMW with **end-to-end ownership** from ingestion through deployment, delivering 80% accuracy improvement.
- Designed **structured observability** and **automated anomaly detection** across healthcare, mobility, and industrial domains -- systematic post-deployment iteration at enterprise scale.
- Translated ambiguous enterprise problems into automation solutions with clear success criteria and evaluation methodologies across multiple clients.

AAM IT GmbH / Vorwerk Switzerland Jan – Oct 2018

Software Developer (Contract) | Munich, Germany

- Engineered backend APIs and data migration pipelines for an enterprise platform; built Python automation for XML validation, anomaly detection, and ETL optimisation.

SELECTED PROJECTS

Multi-Agent Productivity Assistant Gen AI Academy APAC 2026

Production agentic system: orchestrator routing intent to four specialist sub-agents via **multi-step reasoning and tool use**, parallel execution, real-time MCP integration, and AlloyDB AI persistent memory. Cloud Build CI/CD and Cloud Run. Serving real users.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Run · AlloyDB AI

WorkflowEnv – LLM Evaluation & Scoring System Meta x HF Hackathon 2026

OpenEnv-compliant **RL evaluation environment**: 7-criterion partial-credit grader, automated scoring pipeline, and telemetry logging. Defined success criteria and evaluation methodology end-to-end. Avg score 0.96/1.0.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · automated scoring · telemetry

TravelAI – End-to-End RAG Pipeline Kaggle Capstone 2025

End-to-end RAG: heterogeneous ingestion, chunking strategies, hybrid vector search, LLM grounding, and systematic retrieval evaluation experiments.

RAG · Pinecone · Embeddings · Semantic search · LLM grounding · Python

EDUCATION

M.Sc. Information & Communication Systems — TU Chemnitz, Germany 2017 – 2021

IEEE-published thesis

B.E. Electronics & Communication Engineering — VTU 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers 2025

MIT-IDSS — Data Science & Machine Learning 2024 – 2025

Kaggle — Generative AI & AI Agents Intensive 2025

SKILLS

Agents & Orchestration: LangChain, LangGraph, Google ADK, multi-agent systems, tool use, ReAct, prompt engineering, guardrails, failure handling

LLM APIs & Models: Gemini 2.5/2.0, OpenAI API, Vertex AI, LLM evaluation, fine-tuning, RAG pipelines, vector databases

Cloud / Infra: GCP, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, AlloyDB AI, Pinecone

Observability: Datadog, MLflow, structured logging, LLMOps, experiment traceability

Languages: Python, FastAPI, Node.js, SQL, Bash