

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda |

Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer specialising in production LLM systems, RAG pipelines, and hybrid LLM + statistical ML approaches for high-stakes decision workflows. Proven impact in regulated environments: 80% accuracy improvement in ML-driven document processing for automotive-sector clients, 70% faster workflows via conversational multi-agent AI with human-in-the-loop design. Strong command of embeddings, vector search, and semantic retrieval — combined with 3+ years operating in regulated enterprise environments (healthcare, finance, industrial). Experienced with GCP/Vertex AI, LangChain, MLflow, Docker, and production observability tooling.

CORE COMPETENCIES

Production LLM Systems

RAG & Embeddings-Based Retrieval

Hybrid LLM + Statistical ML

Human-in-the-Loop Workflows

Vector Search & Semantic Indexing

LLMOps & Model Observability

Python & FastAPI

Regulated Environment ML

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Automated complex document processing (SAP invoice generation) for BMW service centres using ML-driven data validation — delivering **80% accuracy improvement** and 50% processing time reduction in a regulated, high-volume production environment.
- Built hybrid data validation pipelines combining rule-based logic with statistical ML models for structured and semi-structured document extraction — a direct analogue to insurance claims synthesis.
- Designed normalisation rules and automated quality checks for heterogeneous datasets across healthcare, mobility, and industrial domains.
- Strengthened system observability via structured error handling and logging, reducing recurring production failures — applies directly to LLM reliability engineering.

AAM IT GmbH 2018

Software Developer (Contract) · Munich, Germany

- Built Python scripts for real-time data mining, extraction, and correction against continuously-updated production datasets — experience transferable to LLM-driven information extraction.
- Integrated SOAP web services and automated ETL pipelines via Jenkins and Bash; developed C# APIs for complex multi-interface data flows.

SELECTED AI/ML PROJECTS

AI Invoice Agent — Human-in-the-Loop Multi-Agent System Gen AI Academy APAC · 2025

Production conversational AI agent with human-in-the-loop design for invoice creation, payment tracking, and financial analytics. Multi-agent orchestration with Gemini 2.5 Flash; resolved critical LLM reliability issues in production. **Impact: ~70% faster workflows.**

Gemini 2.5 · LangChain · Multi-agent orchestration · Docker · Cloud Run · HITL design

TravelAI — RAG Pipeline with Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG implementation: data ingestion, embedding-based retrieval indexing, hybrid semantic search, LLM grounding, and controlled experiments measuring retrieval relevance and consistency — hybrid LLM + retrieval approach directly applicable to medical document synthesis.

RAG · Pinecone-style vector DB · Semantic search · Embeddings · LLM grounding · Python

AI SRE Copilot — LLM Observability & Reliability DEVPOST Hackathon · 2025

Multimodal AI copilot for automated root-cause analysis using Vertex AI + Datadog. Demonstrates production LLM monitoring, structured output design, and reliability engineering.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMOps

Generative AI Search Agent Personal Project · Ongoing

LLM-powered semantic retrieval agent: embeddings + vector DB, prompt chaining, contextual memory. **~40% faster resolution vs. keyword search.** Demonstrates hybrid retrieval + LLM approach.

LangChain · Embeddings · Vector DB · OpenAI API · Docker · CI/CD

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial data processing with ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Semantic Search)2025
- MIT-IDSS — Data Science & Machine Learning · HarvardX CS50 Python · Kaggle GenAI & AI Agents Intensive2024–2025

SKILLS

LLMs & GenAI: OpenAI API, Gemini 2.5, LangChain, LangGraph, RAG, embeddings, vector DBs (Pinecone), prompt engineering, LLMOps

ML Engineering: Supervised/unsupervised ML, statistical models, MLflow, evaluation design, feature engineering

Cloud / MLOps: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, Datadog observability **Languages:** Python, FastAPI, Node.js, SQL, C#, Bash