

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer with production experience building high-concurrency agentic AI systems and LLM-powered microservices — directly applicable to Fastino's agentic AI platform. Shipped two multi-agent systems deployed on GCP Cloud Run handling real user workloads, a FastAPI-based RL evaluation environment validated at 0.96/1.0, and a conversational AI agent with full-stack production ownership. Comfortable across the stack from LLM API integration and prompt engineering to FastAPI microservice design, vector DB tuning, and Docker deployment. IEEE-published ML research and 3 years enterprise backend engineering adds the systems rigour that high-performance AI infrastructure demands. Open to joining a fast-moving startup building the next generation of efficient LLM systems.

CORE COMPETENCIES

Production Agentic AI Systems

LLM API Integration (Gemini, OpenAI)

FastAPI Microservice Design

Vector DB Tuning & Semantic Retrieval

Multi-Agent Orchestration (Google ADK)

Docker & GCP Cloud Run Deployment

LLM Evaluation & Prompt Engineering

LangChain & LangGraph

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Designed and deployed ML-driven backend systems for BMW service centres — full pipeline from data validation and model architecture through production deployment, delivering **80% accuracy improvement** at high volume.
- Built backend systems across healthcare, mobility, and industrial domains; implemented structured observability and automated anomaly detection to maintain production reliability at scale.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — High-Concurrency Agentic System Gen AI Academy APAC · 2026

Architected and shipped a production multi-agent system handling real user workloads: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel agent execution resolving multi-step requests concurrently, real-time MCP tool-use, AlloyDB AI persistent memory. Deployed on Cloud Run. Demonstrates the high-concurrency agentic workload design Fastino's platform requires.

Gemini 2.0 Flash · Google ADK · LangGraph · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

WorkflowEnv — LLM Agent Evaluation Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed a FastAPI-based RL evaluation environment for LLM workflow agents — 7-criterion partial-credit grader, telemetry logging, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0, 3/3 tasks passing.** Demonstrates systematic LLM evaluation and API-first microservice design at production quality.

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM evaluation · telemetry

AI Invoice Agent — Production LLM Microservice Gen AI Academy APAC · 2026

Designed and shipped a conversational multi-agent AI with FastAPI + Node.js backend in production — Gemini 2.5 Flash powering invoice creation, payment tracking, and analytics. Diagnosed and resolved a critical production LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.**

Gemini 2.5 Flash · LangChain · FastAPI · Node.js · Docker · Cloud Run

TravelAI — RAG Pipeline with Vector DB Tuning Kaggle Capstone · 2025

End-to-end RAG system with controlled experiments measuring relevance across chunking strategies, embedding models, and retrieval configurations. Demonstrates vector DB tuning and LLM grounding for information extraction — applicable to Fastino's task-specific extraction and classification systems.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Python

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

LLM & Agents: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, multi-agent orchestration, prompt engineering, LLM evaluation, RAG pipelines

AI Systems: Agentic workflows, vector DBs (Pinecone), embeddings, semantic search, information extraction, RL environment design, OpenEnv

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, MLflow, Datadog

Languages: Python, FastAPI, Node.js, TypeScript, SQL, Bash