

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer specialising in building evaluation infrastructure that ensures LLM agents perform reliably at enterprise scale — the precise scope of Fieldguide's AI Engineer, Quality role. Built an OpenEnv-compliant evaluation environment from scratch: 7-criterion partial-credit grader, automated scoring pipeline, strict telemetry logging for full observability — validated at 0.96/1.0 across 3/3 tasks (Meta × Hugging Face 2026). Shipped two production multi-agent systems on GCP Cloud Run with real users and a RAG system with controlled retrieval experiments. Strong command of LLM evaluation design, production agent reliability, and the observability tooling that enterprise clients of top-100 accounting firms require to trust AI-powered audit workflows.

CORE COMPETENCIES

LLM Agent Evaluation Infrastructure

Automated Scoring & Telemetry Pipelines

Production Agent Reliability Engineering

LangChain & LangGraph

Multi-Agent Orchestration (Google ADK)

RAG Pipelines & Retrieval Evaluation

FastAPI & Docker CI-CD

Observability & MLflow / Datadog

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven systems for BMW service centres — owning pipelines from model evaluation and validation through production deployment and monitoring, delivering **80% accuracy improvement**.
- Implemented structured observability, automated failure detection, and root-cause analysis across healthcare, mobility, and industrial domains — the reliability discipline that enterprise AI audit workflows demand.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — LLM Agent Evaluation Infrastructure OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment for LLM workflow agents — exactly the evaluation infrastructure Fieldguide requires for AI agents at enterprise scale. Built a 7-criterion partial-credit grader covering task completion, instruction-following, output grounding, constraint adherence, and failure handling. Strict telemetry logging ensures full audit trails for every evaluation run. Automated scoring pipeline, Docker deployment to HF Spaces. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL evaluation design · LLM agent evaluation · telemetry logging · automated scoring

Multi-Agent Productivity Assistant — Production Agent System Gen AI Academy APAC · 2026

Architected and shipped a multi-agent system from scratch: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel execution, real-time MCP tool-use, AlloyDB AI persistent memory. Deployed on Cloud Run with real users. Production deployment experience informs evaluation work — you understand failure modes because you’ve hit them.

Gemini 2.0 Flash · Google ADK · LangGraph · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

TravelAI — RAG Pipeline with Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG system with controlled experiments measuring retrieval relevance and consistency across chunking strategies, embedding models, and retrieval configurations — systematic evaluation methodology applicable to grounding and factual reliability in audit-domain LLM agents.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

AI SRE Copilot — LLMops Observability DEVPOST Hackathon · 2025

LLMops observability infrastructure integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis of production incidents. FastAPI on Cloud Run. Demonstrates real-time LLM monitoring and structured AI diagnostics — the observability tooling that underpins reliable AI agent quality at scale.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMops

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMops, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

- Eval & Observability: LLM evaluation pipeline design, partial-credit grading, automated scoring, telemetry logging, guardrails, hallucination detection, MLflow, Datadog, structured observability
- LLMs & Agents: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, multi-agent orchestration, RAG pipelines, vector DBs (Pinecone), prompt engineering
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, AlloyDB AI Languages: Python, FastAPI, Node.js, SQL, Bash