

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer specialising in LLM-powered agent workflows, production observability, and cloud-native AI systems on GCP. Built and shipped a multimodal AI SRE Copilot integrating Datadog with Vertex AI for automated root-cause analysis, and an OpenEnv-compliant RL evaluation environment for LLM workflow agents validated at 0.96/1.0 average score. Deep experience with multi-agent orchestration (Google ADK, LangGraph), prompt engineering, and LLMops monitoring — directly applicable to Grafana's incident lifecycle automation mission. Proven production impact: 80% accuracy improvement in ML-driven enterprise automation; 70% faster workflows via agentic AI. Comfortable owning end-to-end AI systems from data pipelines through deployment and observability on Cloud Run.

CORE COMPETENCIES

LLM-Powered Agent Workflows

Production Observability & LLMops

Multi-Agent Orchestration (Google ADK)

GCP / Vertex AI / Cloud Run

Datadog Integration & Monitoring

Incident Automation & Root-Cause Analysis

Prompt Engineering & LLM Evaluation

FastAPI & Docker / CI-CD

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML systems for BMW service centres — from data pipelines and model design through production deployment and monitoring — delivering **80% accuracy improvement** in automated invoice generation.
- Strengthened production observability via structured error handling, logging, and automated failure detection — reducing recurring failures through systematic root-cause analysis; directly transferable to LLM operational workflows.
- Built backend data systems across healthcare, mobility, and industrial domains in regulated, high-volume production environments.

AAM IT GmbH / Vorwerk Switzerland 2018

Software Developer (Contract) · Munich, Germany

- Built automated ETL pipelines and C# APIs for multi-interface integrations; cross-timezone production support ensuring pipeline stability.

SELECTED AI/ML PROJECTS

AI SRE Copilot — LLM-Powered Observability & Incident Automation DEVPOST Hackathon · 2025

Built multimodal AI system integrating Datadog metrics and logs with Vertex AI (Gemini) for automated root-cause analysis of production incidents. FastAPI service deployed on Cloud Run delivers structured AI insights and real-time performance signals — directly models the LLM-powered Grafana Ops workflows the role requires.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMops observability · Python

WorkflowEnv — OpenEnv RL Environment for LLM Agents OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents — FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, and strict telemetry logging. Both Phase 1 and Phase 2 validated. Avg score 0.96/1.0, 3/3 tasks passing. Demonstrates agent evaluation infrastructure at production quality.

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry logging

Multi-Agent Productivity Assistant — Orchestrated LLM Workflows Gen AI Academy APAC · 2026

Architected multi-agent system with Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK. Real-time tool-use, parallel agent execution for multi-step workflows, AlloyDB AI persistent memory. Deployed on Cloud Run.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Production LLM Reliability Engineering Gen AI Academy APAC · 2026

Diagnosed and resolved a critical production LLM reliability issue (model/API version mismatch) in a live multi-agent system — restoring stable low-latency responses. Impact: ~70% faster workflows. Demonstrates the operational reliability mindset Grafana's AI systems require.

Gemini 2.5 Flash · LangChain · Docker · Cloud Run · LLM reliability engineering

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale ML pipeline for geospatial data processing adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

LLM & Agent: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, prompt engineering, LLM evaluation, multi-agent orchestration, ReAct pattern

Observability: Datadog (metrics, logs, APM), Vertex AI, structured AI insights, LLMOps monitoring, MLflow

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces Languages: Python, FastAPI, Node.js, SQL, Bash