

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer who turns manual, spreadsheet-driven business workflows into LLM-powered automated systems — exactly what HappyCo's Business Transformation role requires. Built two production agentic AI systems deployed on GCP Cloud Run: a multi-agent assistant automating Calendar, Task, Notes, and Email workflows (~70% faster operations), and a conversational AI agent replacing manual invoice processing. Three years of enterprise backend engineering across BMW, healthcare, and industrial domains, delivering measurable efficiency gains (80% accuracy improvement, 40% workload reduction). Comfortable working directly with business stakeholders to identify inefficiencies and design Python + LLM API solutions that fit into real workflows.

CORE COMPETENCIES

AI-Driven Business Transformation

Workflow Automation & LLM APIs

Multi-Agent Productivity Systems

Python & FastAPI

Production AI Deployment (Cloud Run)

Business Process Automation

LangChain & Google ADK

End-to-End AI Delivery

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Automated business-critical ML workflows for BMW service centres — replacing manual, spreadsheet-driven processes with production systems, delivering **80% accuracy improvement** in automated invoice generation and measurable operational efficiency gains.
- Built backend automation systems across healthcare, mobility, and industrial domains; worked directly with domain experts to identify inefficiencies and design solutions that integrated into existing business workflows.
- Implemented observability and automated failure detection, reducing recurring production failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

AI Invoice Agent — Business Workflow Automation Gen AI Academy APAC · 2026

Designed and shipped a conversational AI agent automating invoice creation, payment tracking, and real-time analytics — replacing manual workflows end-to-end. Diagnosed and resolved a critical production LLM reliability issue mid-flight. **Impact: ~70% faster workflows.** Full-stack ownership: React + Vite frontend, Node.js backend, Docker, Cloud Run. The kind of internal business AI tool this role requires.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

Multi-Agent Productivity Assistant — Ops Workflow Automation Gen AI Academy APAC · 2026

Built from zero to production: Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK. Real-time tool-use via Google Calendar MCP and Gmail MCP. Parallel agent execution resolves multi-step user requests in a single turn. AlloyDB AI persistent memory. Deployed on Cloud Run with real users — directly models the productivity automation HappyCo needs.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI SRE Copilot — LLM-Powered Operations Automation DEVPOST Hackathon · 2025

Multimodal AI application integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis of production incidents. FastAPI service on Cloud Run delivering structured AI insights — demonstrates LLM API integration for operational workflows.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run

WorkflowEnv — LLM Agent Evaluation OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents — 7-criterion partial-credit grader, telemetry logging, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.**

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

- LLM & Agents: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, multi-agent orchestration, prompt engineering, workflow automation
- AI Systems: RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM evaluation, business process automation
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, Datadog, MLflow
- Languages: Python, FastAPI, Node.js, TypeScript, React, SQL, Bash, C#