Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

ML Engineer with hands-on experience designing and deploying LLM evaluation frameworks and RL-based training environments — bridging research and engineering in exactly the way Innodata's LLM Post-Training & Evaluation role demands. Built an OpenEnv-compliant RL evaluation environment from scratch: 7-criterion partial-credit grader, telemetry logging for full experiment traceability, automated scoring pipeline — validated at 0.96/1.0 across all tasks (Meta × Hugging Face 2026). Practical command of LLM APIs, RAG pipelines, agent orchestration, and Python ML infrastructure, grounded by three years of enterprise software engineering in regulated domains. Ready to design the pipelines and tooling that connect data, evaluation, and post-training for foundation model builders.

CORE COMPETENCIES

LLM Evaluation Framework Design | RL Environment & Post-Training Pipelines | Automated Scoring & Experiment Traceability | Preference Optimisation Concepts (RLHF/DPO) | RAG Pipelines & Data Integration | LangChain & LangGraph | Python ML Infrastructure | MLflow & Datadog Observability

WORK EXPERIENCE

PICA GmbH
2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven data validation and automation pipelines for BMW service centres — end-to-end ownership from feature engineering, model evaluation, and experiment tracking through production deployment, delivering **80% accuracy improvement**.
- Implemented structured observability and automated failure detection across healthcare, mobility, and industrial domains — the experiment reliability and production engineering discipline Innodata's model improvement work requires.

AAM IT GmbH / Vorwerk Switzerland
Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes
May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — RL Evaluation Environment for LLM Post-Training
OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment — the exact type of tooling Innodata's post-training pipeline work requires. Built a 7-criterion partial-credit grader integrating evaluation signals into an RL feedback loop: task completion, instruction adherence, constraint satisfaction, output quality, and scoring confidence. Full telemetry logging for experiment traceability and audit trails. Automated scoring pipeline connecting agent outputs to reward signals. Docker deployment to HF Spaces. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM evaluation · telemetry · automated scoring · reward modelling

TravelAI — RAG Pipeline with Data Integration & Evaluation   Kaggle Capstone · 2025

End-to-end RAG system: heterogeneous data ingestion, chunking and embedding pipeline design, hybrid vector search, LLM grounding, and controlled experiments measuring retrieval relevance and response consistency across embedding models and retrieval configurations. Demonstrates systematic data integration and LLM component evaluation — directly applicable to Innodata's data-to-training pipeline work.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python · FastAPI

Multi-Agent Productivity Assistant — Production LLM System   Gen AI Academy APAC · 2026

Architected multi-agent AI from scratch: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel execution, real-time MCP tool-use, AlloyDB AI persistent memory. Deployed on Cloud Run with real users. Production experience with LLM behaviour under real workloads informs evaluation design — understanding where models fail drives better harnesses.

Gemini 2.0 Flash · Google ADK · LangGraph · LangChain · FastAPI · Cloud Run · AlloyDB AI

AI SRE Copilot — LLMOps & Monitoring Infrastructure   DEVPOST Hackathon · 2025

LLMOps observability infrastructure integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis. FastAPI on Cloud Run. Demonstrates structured logging, MLflow-style experiment tracking concepts, and production LLM monitoring — the observability backbone for continuous model improvement loops.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMOps · MLflow

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany   2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University   2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)   2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive   2024–2025

SKILLS

LLM Eval & Training: RL environment design, LLM evaluation frameworks, partial-credit grading, automated scoring, RLHF/DPO concepts, SFT pipeline design, preference optimisation, telemetry, experiment traceability

ML & LLMs: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, prompt engineering

MLOps / Infra: MLflow, Datadog, structured observability, GCP, Vertex AI, Cloud Run, Docker, FastAPI, Cloud Build CI/CD, HF Spaces

Languages: Python, FastAPI, SQL, Bash, Node.js