

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer specialising in stateful multi-agent architecture and production LLM systems — the specific skills Instrument's simulation engine contract requires. Built a production multi-agent system using **Google ADK** (explicitly preferred by this role) with parallel agent execution, persistent AlloyDB AI memory across thousands of sessions, and real-time MCP tool-use — deployed on Cloud Run with real users. Designed an OpenEnv RL evaluation environment with 7-criterion partial-credit grading and telemetry logging, validated at 0.96/1.0. Hands-on experience with LangGraph, LangChain, context optimisation, and prompt engineering for complex multi-step agent workflows. Available for immediate W-2 contract engagement.

CORE COMPETENCIES

Multi-Agent Architecture (Google ADK)

Stateful Agent Design & Session Memory

LangGraph & Agent Orchestration

Context & Token Optimisation

Production LLM Systems

Agent Evaluation & Telemetry

Prompt Engineering & Guardrails

FastAPI & GCP Cloud Run

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML systems for BMW service centres — from data pipelines and model design through production deployment and monitoring — delivering **80% accuracy improvement** in automated invoice generation.
- Built production backend systems across healthcare, mobility, and industrial domains; implemented structured observability and automated failure detection, reducing recurring production failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Google ADK Production System Gen AI Academy APAC · 2026

Architected a stateful multi-agent system using **Google ADK** (the framework Instrument specifically requires): Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) with parallel execution resolving multi-step requests in a single turn. Real-time tool-use via Google Calendar MCP and Gmail MCP. AlloyDB AI persistent memory maintains context across sessions. Deployed on Cloud Run with real users — directly models the stateful simulation engine architecture this contract requires.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI · MCP

WorkflowEnv — Agent Evaluation with Telemetry & Grading

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL environment for LLM workflow agents — 7-criterion partial-credit grader, strict telemetry logging for full observability, Docker deployment to HF Spaces. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.** Demonstrates the evaluation infrastructure and guardrail design relevant to complex multi-agent simulation systems.

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry

AI Invoice Agent — Production LLM Reliability

Gen AI Academy APAC · 2026

Shipped conversational multi-agent AI from prototype to production; diagnosed and resolved a critical LLM reliability issue (model/API version mismatch) in a live system mid-flight. **Impact: ~70% faster workflows.** Demonstrates the operational reliability mindset required for high-volume agent systems.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

AI SRE Copilot — Multimodal LLM Application

DEVPOST Hackathon · 2025

Multimodal AI application integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis. FastAPI on Cloud Run. Demonstrates context-rich LLM pipeline design for complex, data-dense inputs — relevant to managing global system state in large-scale simulations.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

- Multi-Agent: Google ADK, LangGraph, LangChain, multi-agent orchestration, stateful agent design, MCP tool-use, parallel execution, ReAct pattern, prompt engineering
- LLMs & Eval: Gemini 2.5/2.0, OpenAI API, context optimisation, token management, LLM evaluation, guardrails, telemetry logging, RL environment design
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, AlloyDB AI, HF Spaces
- Languages: Python, FastAPI, Node.js, TypeScript, SQL, Bash