

# Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda |

Bangalore, India · Open to relocation, Zurich Switzerland

## PROFESSIONAL SUMMARY

AI/ML Engineer with production experience delivering RAG implementations, AI agent systems, and LLM-powered applications using Gemini, OpenAI, and LangChain. Proven track record: 80% accuracy improvement in ML-driven finance automation for BMW-sector clients, 70% faster invoice workflows via conversational multi-agent AI. Three years of professional experience in Germany (Munich) and M.Sc. from TU Chemnitz — a combination that positions me to contribute directly to KPMG Switzerland's AI practice with both the technical depth and the European professional context clients expect. Experienced across GCP/Vertex AI, prompt engineering, model evaluation, and MLOps best practices in production environments.

## CORE COMPETENCIES

RAG Pipeline Design

LLM Agents & Orchestration

Gemini 2.5 / OpenAI API / Mistral

Prompt Engineering & Model Evaluation

Python & FastAPI

MLOps Best Practices

GCP / Vertex AI / Cloud Run

Multi-Agent Systems (Google ADK)

## WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Delivered **80% accuracy improvement** in ML-driven SAP invoice automation for BMW service centres across South Germany and Austria — applying model evaluation, validation logic, and automated data quality checks in a regulated production environment.
- Built backend data systems for healthcare, mobility, and industrial domains; designed normalisation rules and anomaly detection pipelines for structured and semi-structured datasets.
- Strengthened observability via structured error handling and logging, reducing recurring production failures — directly applicable to LLMOps monitoring practices.
- Collaborated cross-functionally with business stakeholders across German-language enterprise clients to translate requirements into production ML solutions.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

ADREMES GmbH 2017

Software Developer (Trainee) · Hamburg, Germany

- Backend validation and session-handling workflows; dashboard enhancements for cloud ad-analytics platform.

## SELECTED AI/ML PROJECTS

AI Invoice Agent — Conversational Multi-Agent System

Gen AI Academy APAC · 2026

Built production-grade conversational AI agent automating invoice creation and real-time financial analytics using multi-agent orchestration and Gemini 2.5 Flash. Resolved LLM reliability issues (model/API version mismatch) in production. **Impact: ~70% faster invoice workflows.**

Gemini 2.5 Flash · LangChain · Google ADK · FastAPI · Docker · Cloud Run · Cloud Build CI/CD

Multi-Agent Productivity Assistant — Orchestrated AI System

Gen AI Academy APAC · 2026

Architected an AI agent system where an Orchestrator Agent routes intent to four specialist sub-agents (Calendar, Task, Notes, Email) — applying multi-agent orchestration, real-time tool-use (Google Calendar MCP, Gmail MCP), and stateful memory (AlloyDB AI).

Gemini 2.0 Flash · Google ADK · LangGraph · Node.js · Cloud Run · AlloyDB AI

WorkflowEnv — OpenEnv RL Environment for LLM Agents

OpenEnv AI Hackathon · Meta x Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents. Implemented FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, and strict telemetry logging. Phase 1 and Phase 2 both validated — **avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry logging

TravelAI — RAG Concierge

Kaggle Capstone · 2025

End-to-end RAG pipeline: heterogeneous data ingestion, embedding-based retrieval indexes, semantic search, LLM grounding, and controlled retrieval experiments measuring relevance and consistency. Directly demonstrates RAG implementation skills.

RAG · Pinecone-style Vector DB · Semantic Search · LLM Grounding · Python

AI SRE Copilot — Multimodal LLM Observability

DEVPOST Hackathon · 2025

Multimodal AI copilot for automated root-cause analysis integrating Datadog metrics/logs with Vertex AI (Gemini). FastAPI service on Cloud Run — demonstrates LLMOps observability and production monitoring.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · MLOps

EDUCATION

|                                                                                                                               |             |
|-------------------------------------------------------------------------------------------------------------------------------|-------------|
| M.Sc. Information & Communication Systems · TU Chemnitz, Germany                                                              | 2017 – 2021 |
| Thesis: IEEE-published — large-scale geospatial data processing adopted by a Galileo Map Service Provider for production use. |             |
| B.E. Electronics & Communication Engineering · Visvesvaraya Technological University, India                                   | 2013 – 2017 |

CERTIFICATIONS

|                                                                                                                        |           |
|------------------------------------------------------------------------------------------------------------------------|-----------|
| DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Semantic Search, Embeddings) | 2025      |
| DataCamp — Developing AI Applications · OpenAI Fundamentals                                                            | 2025      |
| MIT-IDSS — Data Science & Machine Learning · HarvardX CS50 Python · Kaggle Generative AI & AI Agents Intensive         | 2024–2025 |

SKILLS

**LLMs & GenAI:** Gemini 2.5, OpenAI API, Mistral (via API), LangChain, LangGraph, RAG, prompt engineering, fine-tuning, LLMOps

**Agentic AI:** Multi-agent orchestration, Google ADK, tool-use, stateful agent design, ReAct pattern

**Cloud / MLOps:** GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, Datadog, MLflow    **Languages:** Python, TypeScript/Node.js, FastAPI, SQL, C#, Bash

**Languages spoken:** English (fluent), German (professional — 3 years work experience in Germany)