

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer specialising in LLM evaluation, alignment, and guardrail design for production systems — the precise scope of LeoTech's AI/LLM Evaluation & Alignment role. Built and deployed an OpenEnv-compliant LLM evaluation framework from scratch: 7-criterion partial-credit grader, telemetry logging, automated scoring pipeline, validated at 0.96/1.0 (Meta × Hugging Face 2026). Hands-on with LangChain, LangGraph, and Python evaluation pipelines — the exact stack LeoTech requires. Two production multi-agent systems on GCP Cloud Run demonstrate the practical AI deployment experience that informs rigorous safety and alignment work. Three years of enterprise software engineering in high-stakes domains (healthcare, automotive, mobility) adds the reliability discipline that public safety AI systems demand.

CORE COMPETENCIES

LLM Evaluation Framework Design

AI Alignment & Guardrail Engineering

LangChain & LangGraph

Automated Scoring & Telemetry Pipelines

Production LLM Systems

Bias Detection & Mitigation

RAG Pipelines & Semantic Retrieval

Multi-Agent Orchestration (Google ADK)

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven data validation and automation systems for BMW service centres — end-to-end pipeline ownership from feature engineering and model evaluation through deployment and monitoring, delivering **80% accuracy improvement**.
- Operated in high-stakes, regulated domains (healthcare, mobility, industrial); implemented structured observability and automated failure detection, applying the rigour and reliability standards that safety-critical AI systems require.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — LLM Evaluation & Alignment Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment for LLM workflow agents — built to assess model behaviour rigorously and safely. 7-criterion partial-credit grader covering task completion, instruction-following, output quality, and constraint adherence. Full telemetry logging for evaluation transparency and audit trails. Docker deployment to HF Spaces. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.** Directly demonstrates the LLM evaluation pipeline design, alignment-aware grading, and observability infrastructure LeoTech's role requires.

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM evaluation · telemetry logging · automated scoring

TravelAI — RAG Pipeline with Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG system with controlled experiments measuring retrieval relevance, consistency, and factual grounding across chunking strategies, embedding models, and retrieval configurations. Demonstrates quantitative LLM component evaluation and systematic analysis of model output quality — methods applicable to alignment and bias auditing.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

Multi-Agent Productivity Assistant — Production LLM System Gen AI Academy APAC · 2026

Built and shipped a multi-agent AI system handling real user workloads: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel execution, AlloyDB AI persistent memory, real-time MCP tool-use. Deployed on Cloud Run. Practical production experience informs evaluation work — you understand what you're testing.

Gemini 2.0 Flash · Google ADK · LangGraph · LangChain · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Production LLM Reliability Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI from prototype to production; diagnosed and resolved a critical LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Demonstrates the debugging and root-cause mindset needed to identify model misbehaviour in live systems.

Gemini 2.5 Flash · LangChain · FastAPI · Node.js · Docker · Cloud Run

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

LLM Eval & Alignment: LLM evaluation pipeline design, partial-credit grading, guardrails, bias detection, telemetry logging, automated scoring, RLHF concepts, OpenEnv

LLMs & Agents: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, multi-agent orchestration, prompt engineering, RAG pipelines, vector DBs (Pinecone)

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, MLflow, Datadog

Languages: Python, FastAPI, Node.js, SQL, Bash