

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

Backend AI Engineer with a track record of shipping production LLM-powered systems — the core of Mintlify's AI feature work. Built two agentic AI systems deployed on GCP Cloud Run with real users, a RAG pipeline with controlled retrieval experiments, and a multimodal AI observability application — all on FastAPI backends with Docker CI/CD. Strong command of LLM API integration (Gemini, OpenAI), multi-agent orchestration, production reliability engineering, and full-stack delivery from prototype to deployment. Three years of enterprise backend engineering adds the systems discipline that developer-facing AI infrastructure demands. Genuinely excited to build AI features for the toolchain that Anthropic, Cursor, and Scale AI use.

CORE COMPETENCIES

Backend AI Features & LLM Integration

Production LLM Reliability Engineering

FastAPI & Python Backend Systems

RAG Pipelines & Semantic Retrieval

Multi-Agent Orchestration (Google ADK)

Docker & GCP Cloud Run CI-CD

LLM API Integration (Gemini, OpenAI)

LangChain & LangGraph

WORK EXPERIENCE

PICA GmbH

2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML backend systems for BMW service centres — from data pipelines and model integration through production deployment and monitoring — delivering **80% accuracy improvement**.
- Built backend systems across healthcare, mobility, and industrial domains in regulated environments; implemented structured observability and automated failure detection to maintain production reliability.

AAM IT GmbH / Vorwerk Switzerland

Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes

May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

AI Invoice Agent — Production LLM Backend

Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI from prototype to production on a FastAPI + Node.js backend — Gemini 2.5 Flash powering invoice creation, payment tracking, and analytics. Diagnosed and resolved a critical production LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Full-stack: React + Vite frontend, Docker, Cloud Run. Demonstrates the backend LLM reliability engineering Mintlify's AI features require.

Gemini 2.5 Flash · LangChain · FastAPI · React + Vite · Node.js · Docker · Cloud Run

Multi-Agent Productivity Assistant — Production Agentic System

Gen AI Academy APAC · 2026

Architected and shipped a multi-agent system from scratch on FastAPI + Cloud Run: Orchestrator routing to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, real-time MCP tool-use, parallel execution, AlloyDB AI persistent memory. Real users in production.

Gemini 2.0 Flash · Google ADK · LangGraph · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

TravelAI — RAG Backend with Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG backend: ingestion pipelines, chunking and embedding policies, vector search, LLM grounding, and controlled experiments measuring retrieval relevance across configurations. Demonstrates the semantic search and knowledge retrieval backend work relevant to Mintlify’s AI-powered documentation features.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Python · FastAPI

WorkflowEnv — LLM Agent Evaluation Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow agents — FastAPI backend, 7-criterion partial-credit grader, Docker deployment to HF Spaces, telemetry logging. Both phases validated. **Avg score 0.96/1.0.**

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · telemetry

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

Backend & LLM: Python, FastAPI, Node.js, LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, multi-agent orchestration, prompt engineering

AI Systems: RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM evaluation, production LLM reliability, ReAct pattern

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, HF Spaces, Datadog, MLflow

Languages: Python, FastAPI, Node.js, TypeScript, React, SQL, Bash