

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer with production experience building agentic AI platforms, RAG systems, and LLMOps pipelines for enterprise workloads. Strong command of LangChain, multi-agent orchestration, vector databases, and compliance-first AI design — applied in regulated environments including healthcare and financial services. Proven impact: 80% accuracy improvement in ML-driven data validation for automotive-sector clients; 70% faster workflows via production multi-agent AI. Experienced with GCP/Vertex AI, Datadog observability, MLflow experiment tracking, and containerised CI/CD deployments. M.Sc. in Information Systems (TU Chemnitz, Germany) with an IEEE-published ML thesis adopted by a production Galileo Map Service Provider.

CORE COMPETENCIES

Agentic AI Architecture (LangChain, LangGraph)

RAG at Scale & Vector DB Design

LLMOps & Datadog Observability

Multi-Agent Orchestration (Google ADK)

MLflow & Experiment Tracking

Compliance-First AI in Regulated Environments

GCP / Vertex AI / Cloud Run

Docker & Cloud Build CI/CD

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven validation and automation systems for BMW service centres — achieving **80% accuracy improvement** and 50% processing time reduction in a regulated, high-volume production environment (GDPR, data integrity compliance).
- Operated across healthcare, mobility, and industrial domains — directly relevant to HIPAA/CLIA-adjacent compliance requirements in clinical AI platforms.
- Designed normalisation rules, label quality checks, and automated anomaly detection pipelines for structured and semi-structured datasets; implemented observability via structured error handling and logging.

AAM IT GmbH / Vorwerk Switzerland 2018

Software Developer (Contract) · Munich, Germany

- Designed and executed large-scale database migration with evolving schema; built Python scripts for real-time data mining and ETL pipeline optimisation against continuously-updated datasets.
- Developed C# APIs and integrated SOAP web services via Jenkins/Bash; provided cross-timezone support ensuring pipeline stability.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Agent Orchestration Platform

Gen AI Academy APAC · 2026

Architected multi-agent AI system (Gemini 2.0 Flash + Google ADK): Orchestrator Agent routes intent to four specialist sub-agents (Calendar, Task, Notes, Email) via real-time tool-use (Google Calendar MCP, Gmail MCP). Parallel agent execution handles multi-step workflows in a single turn. AlloyDB AI for persistent structured memory. Demonstrates the agentic AI platform architecture Natera requires.

Gemini 2.0 Flash · Google ADK · LangGraph · Multi-agent orchestration · Cloud Run · AlloyDB AI

AI Invoice Agent — Production Agentic System

Gen AI Academy APAC · 2026

Production multi-agent system automating invoice creation, payment tracking, and analytics using Gemini 2.5 Flash. Resolved critical LLM reliability issue in production. Deployed via Docker + Cloud Build CI/CD on Cloud Run. **Impact: ~70% faster invoice workflows.**

Gemini 2.5 · LangChain · Docker · Cloud Run · CI/CD · LLM reliability engineering

WorkflowEnv — OpenEnv RL Environment for LLM Agents

OpenEnv AI Hackathon · Meta x Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents. Implemented FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, and strict telemetry logging. Phase 1 and Phase 2 both validated — **avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry logging

TravelAI — RAG Pipeline at Scale

Kaggle Capstone · 2025

End-to-end RAG implementation: ingestion pipelines, chunking/embedding policies, hybrid vector search, LLM grounding, retrieval relevance experiments. Demonstrates RAG-at-scale engineering across the full LLMOps stack.

RAG · Pinecone-style vector DB · Embeddings · Semantic search · LLM grounding

AI SRE Copilot — LLMOps Observability

DEVPOST Hackathon · 2025

Multimodal AI copilot integrating Datadog metrics/logs with Vertex AI (Gemini) for automated root-cause analysis. FastAPI on Cloud Run. Demonstrates real-time LLMOps monitoring — cost tracking, drift detection, safety monitoring of AI outputs.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMOps

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany	2017 – 2021
Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.	
B.E. Electronics & Communication Engineering · Visvesvaraya Technological University	2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI)	2025
MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive	2024–2025

SKILLS

Agentic AI: LangChain, LangGraph, Google ADK, multi-agent orchestration, tool-use, ReAct pattern, stateful agent design
LLMs & RAG: Gemini 2.5, OpenAI API, RAG pipelines, vector DBs (Pinecone), embeddings, prompt engineering, fine-tuning
MLOps: MLflow, Datadog, GCP Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, containerisation Languages: Python, FastAPI, Node.js, SQL, Bash