

Darshan Linge Gowda

darshanl1711@gmail.com

linkedin.com/in/darshan-lingegowda-ai

github.com/DarshanLingegowda

Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer with deep hands-on experience building production RAG pipelines, semantic search systems, and LLM evaluation frameworks — the core stack of applied ML at Perplexity. Built end-to-end retrieval systems with controlled experiments measuring relevance and consistency across chunking strategies, embedding models, and retrieval configurations. Strong command of Gemini and OpenAI APIs, vector databases, and LLM grounding. Proven production impact: 80% accuracy improvement in ML-driven enterprise systems; 70% faster workflows via conversational multi-agent AI. IEEE-published ML thesis on large-scale geospatial data processing adopted in production by a map service provider.

CORE COMPETENCIES

RAG Pipelines & Semantic Search

LLM API Integration (Gemini, OpenAI)

Embedding Pipelines & Vector DBs

LLM Evaluation & Retrieval Experiments

Production LLM Reliability

Multi-Agent Orchestration

GCP / Vertex AI / Cloud Run

Python & FastAPI

WORK EXPERIENCE

PICA GmbH

2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven information extraction and synthesis systems for BMW service centres — delivering **80% accuracy improvement** in production; applied feature engineering, model evaluation, and data quality pipelines at scale.
- Designed normalisation rules, label quality checks, and automated anomaly detection for structured and semi-structured datasets in regulated, high-volume environments.
- Strengthened production observability via structured error handling and logging, reducing recurring failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland

Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk’s Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes

May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

TravelAI — RAG Concierge with Retrieval Evaluation

Kaggle Capstone · 2025

End-to-end RAG system: heterogeneous data ingestion, chunking and embedding pipeline, vector search, LLM grounding, and controlled retrieval experiments measuring relevance and consistency across retrieval configurations. Demonstrates the retrieval engineering and LLM evaluation skills central to Perplexity’s applied ML work.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

WorkflowEnv — LLM Agent Evaluation Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant evaluation framework for LLM workflow automation agents — 7-criterion partial-credit grader, strict telemetry logging, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.** Demonstrates systematic LLM evaluation design at production quality.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · telemetry logging

Multi-Agent Productivity Assistant — LLM Workflow Design Gen AI Academy APAC · 2026

Orchestrator Agent routing intent to four specialist sub-agents using Gemini 2.0 Flash + Google ADK, with real-time tool-use via MCP servers and parallel agent execution. AlloyDB AI for persistent memory. Deployed on Cloud Run with real users.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Run · AlloyDB AI

AI Invoice Agent — Production LLM Reliability Gen AI Academy APAC · 2026

Shipped conversational multi-agent AI for invoice creation and analytics; resolved a critical production LLM reliability issue (API version mismatch) mid-flight.

Impact: ~70% faster workflows.

Gemini 2.5 Flash · LangChain · React · Node.js · Docker · Cloud Run

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Semantic Search)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

LLMs & RAG: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM grounding, retrieval evaluation

ML Engineering: Feature engineering, supervised/unsupervised ML, MLflow, evaluation design, data pipelines

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI **Languages:** Python, FastAPI, Node.js, SQL, Bash