

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Remote (Worldwide)

PROFESSIONAL SUMMARY

AI Engineer with a track record of taking LLM-powered products from zero to production — building scalable multi-agent systems, LLM evaluation frameworks, and personalized AI-driven features on top of Gemini and OpenAI APIs. Deep understanding of LLM API integration, prompt engineering, and LLM workflow design, applied across conversational agents, RAG concierge systems, and multimodal observability tooling. Built two production agentic systems from scratch in 2025, both deployed on GCP Cloud Run with real users. Comfortable owning end-to-end ML systems across data pipelines, model training, and production integration. Available for async-first worldwide remote collaboration.

CORE COMPETENCIES

LLM API Integration (Gemini, OpenAI)

Prompt Engineering & LLM Evaluation

Zero-to-One Product Delivery

Scalable Multi-Agent LLM Workflows

RAG & Semantic Search

End-to-End ML System Ownership

GCP / Vertex AI / Cloud Run

Production AI Deployment

WORK EXPERIENCE

PICA GmbH

2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML systems for BMW service centres — from data pipelines and model design through production deployment and monitoring — delivering **80% accuracy improvement** in automated invoice generation.
- Built production backend systems across healthcare, mobility, and industrial domains; operated in regulated environments with strict data governance requirements.
- Strengthened production observability via structured error handling and logging; reduced recurring failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk Switzerland

Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk’s Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes

May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — LLM Workflow Design

Gen AI Academy APAC · 2026

Built from zero to production: Orchestrator Agent routes intent to four specialist sub-agents (Calendar, Task, Notes, Email) using Gemini 2.0 Flash + Google ADK. Real-time tool-use via Google Calendar MCP and Gmail MCP. Parallel agent execution resolves multi-step user requests in a single conversation turn. Demonstrates the AI-driven personalized feature development Poe requires.

Gemini 2.0 Flash · Google ADK · LangGraph · Multi-agent LLM workflows · Cloud Run · AlloyDB AI

4/15/26, 10:18 PM

Darshan Linge Gowda — CV

AI Invoice Agent — Conversational AI from Zero to Production

Gen AI Academy APAC · 2026

Designed and shipped a conversational AI agent from scratch — Gemini 2.5 Flash reasoning layer powering invoice creation, payment tracking, and real-time analytics. Resolved LLM reliability issue in production (API version mismatch). **Impact: ~70% faster workflows.** Zero-to-one ownership across the full stack.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

WorkflowEnv — OpenEnv RL Environment for LLM Agents

OpenEnv AI Hackathon · Meta x Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents. Implemented FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, and strict telemetry logging. Phase 1 and Phase 2 both validated — **avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry logging

TravelAI — RAG Concierge with LLM Evaluation

Kaggle Capstone · 2025

End-to-end RAG system with semantic search, context tracking, LLM grounding, and controlled retrieval experiments — measuring relevance and consistency across retrieval configurations. Demonstrates LLM evaluation and scalable LLM workflow design.

RAG · Vector DB · Embeddings · LLM grounding · Retrieval evaluation · Python

AI SRE Copilot — Multimodal LLM Application

DEVPOST Hackathon · 2025

Multimodal AI application integrating Datadog with Vertex AI for root-cause analysis — demonstrates LLM workflow design for complex, data-rich inputs.

FastAPI + Cloud Run deployment.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · Multimodal LLM

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany

2017 – 2021

Thesis: IEEE-published — large-scale ML system for geospatial data processing adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University

2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)

2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive

2024–2025

SKILLS

LLM APIs: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, prompt engineering, LLM evaluation, fine-tuning

AI Systems: Multi-agent orchestration, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, ReAct pattern

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, MLflow, Datadog **Languages:** Python, FastAPI, Node.js, TypeScript, SQL, Bash