

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

Full-Stack AI Software Engineer with hands-on experience building production AI workflow systems — the exact profile Regrello is scaling. Built WorkflowEnv, an OpenEnv-compliant RL evaluation environment for AI workflow agents: automated multi-step task scoring, telemetry logging, and reward pipeline design validated at 0.96/1.0 (Meta × Hugging Face 2026). Shipped a production multi-agent orchestration system with parallel execution and real-time tool-use on Cloud Run, plus a full-stack conversational AI (React + TypeScript + FastAPI) from prototype to production. Strong command of Python, TypeScript/React, FastAPI, LangChain, LangGraph, and GCP — backed by three years of enterprise backend engineering in production-critical workflow-heavy domains.

CORE COMPETENCIES

AI Workflow Automation & Orchestration

Full-Stack Engineering (React/TypeScript + Python)

LangChain & LangGraph

Multi-Agent System Design

FastAPI & Python Backend

RL Evaluation & Automated Scoring

Docker & GCP Cloud Run CI/CD

SQL & Data Pipelines

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven automation pipelines for BMW service centres — end-to-end ownership from workflow design and model integration through production deployment, delivering **80% accuracy improvement** in automated process validation.
- Engineered production backend systems with structured observability across healthcare, mobility, and industrial workflow domains — the reliability and domain breadth Regrello's AI workflow engine requires.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — AI Workflow Agent Evaluation System OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment for AI workflow agents — a direct match for Regrello's AI Engine work. Built a 7-criterion partial-credit grader scoring multi-step workflow task completion, instruction adherence, constraint satisfaction, and output quality. Telemetry logging for experiment traceability. Automated scoring pipeline connecting agent outputs to reward signals. Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL evaluation · workflow agent evaluation · automated scoring · telemetry

Multi-Agent Productivity Assistant — Production AI Workflow Orchestration

Gen AI Academy APAC · 2026

Architected and shipped a production multi-agent workflow system: Orchestrator routing intent across four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, parallel execution resolving complex multi-step requests in a single turn, real-time MCP tool-use, AlloyDB AI persistent memory. Cloud Build CI/CD, Cloud Run deployment. Real users in production — end-to-end workflow automation from natural language to action.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Full-Stack Production AI Application

Gen AI Academy APAC · 2026

Shipped a full-stack conversational AI from prototype to production: React + Vite frontend, FastAPI + Node.js backend, Gemini 2.5 Flash LLM orchestration via LangChain. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Docker + Cloud Run. Full-stack delivery across React/TypeScript + Python — directly applicable to Regrello's stack.

Gemini 2.5 Flash · LangChain · React + Vite · TypeScript · FastAPI · Node.js · Docker · Cloud Run

TravelAI — RAG Pipeline & Data Integration

Kaggle Capstone · 2025

End-to-end RAG system: ingestion pipelines, chunking and embedding, hybrid vector search, LLM grounding, and controlled retrieval experiments across configurations — the data integration and systematic experimentation applicable to Regrello's Large Tabular Model and AI engine data layer.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany

2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University

2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)

2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive

2024–2025

SKILLS

- Full-Stack: Python, FastAPI, React, TypeScript, Node.js, SQL, REST APIs, GraphQL
- AI & Agents: LangChain, LangGraph, Google ADK, multi-agent orchestration, RL evaluation, LLM evaluation, MCP tool-use, RAG pipelines, prompt engineering
- Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, AlloyDB AI
- Data: SQL, data pipelines, vector DBs (Pinecone), embeddings, semantic search, MLflow, Datadog