

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Software Engineer with hands-on experience building production AI systems for data analysis, automated evaluation, and orchestration at scale — the engineering profile Remesh's AI platform requires. Shipped a multi-agent AI system with parallel execution and persistent memory on Cloud Run with real users; an RL evaluation environment with automated scoring and telemetry validated at 0.96/1.0; and an end-to-end RAG pipeline with systematic retrieval experiments. Strong command of Python, LangChain, LangGraph, FastAPI, and GCP infrastructure — backed by three years of enterprise backend engineering in data-heavy, high-volume production environments.

CORE COMPETENCIES

Production AI Systems Engineering

Multi-Agent Orchestration & Parallel Execution

LLM Evaluation & Automated Scoring

RAG Pipelines & Semantic Search

LangChain & LangGraph

FastAPI & Python Backend

GCP Cloud Run & Docker

Data Pipelines & Observability

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven data validation and automation pipelines for BMW service centres — end-to-end ownership from data ingestion and model design through production deployment, delivering **80% accuracy improvement**.
- Implemented structured observability and automated anomaly detection across healthcare, mobility, and industrial domains — the data pipeline reliability and production monitoring discipline Remesh's AI analysis platform requires.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Production AI System at Scale Gen AI Academy APAC · 2026

Architected and shipped a production multi-agent AI system: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel agent execution resolving complex requests in a single turn, real-time MCP tool-use, AlloyDB AI persistent session memory. Cloud Build CI/CD, Cloud Run deployment. Real users in production. Demonstrates production-grade AI orchestration at the scale Remesh's audience analysis platform requires.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

WorkflowEnv — Automated LLM Evaluation & Scoring System OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL evaluation environment — 7-criterion partial-credit grader, telemetry logging for full experiment traceability, automated scoring pipeline connecting agent outputs to feedback signals, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.** Demonstrates systematic AI evaluation methodology applicable to Remesh's response analysis and quality scoring at scale.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · automated scoring · telemetry · experiment traceability

TravelAI — RAG Pipeline with Data Integration & Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG system: heterogeneous data ingestion, chunking and embedding pipeline design, hybrid vector search, LLM grounding, and controlled experiments measuring retrieval relevance and response consistency across configurations — the data integration and LLM grounding discipline relevant to Remesh's AI-powered audience analysis at scale.

RAG · Vector DB · Pinecone · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

AI SRE Copilot — LLMOps Observability Infrastructure DEVPOST Hackathon · 2025

LLMOps observability infrastructure integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis of production incidents. FastAPI on Cloud Run. Demonstrates real-time LLM monitoring, structured logging, and the observability backbone required for reliable AI analysis at scale.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMOps · structured observability

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

AI & Agents: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, multi-agent orchestration, parallel execution, LLM evaluation, automated scoring, RAG pipelines

Data & Retrieval: vector DBs (Pinecone), embeddings, semantic search, retrieval evaluation, data pipelines, LLM grounding, ETL, anomaly detection

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, AlloyDB AI

Observability: Datadog, MLflow, structured logging, LLMOps, experiment traceability · **Languages:** Python, FastAPI, Node.js, SQL, Bash