

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer with production experience building agentic AI platforms, RAG pipelines, and LLMOps infrastructure in regulated environments — directly applicable to Rula's mental health AI mission. Three years operating in GDPR-regulated enterprise environments (healthcare, automotive, industrial) with strict data governance and compliance requirements. Built multi-agent orchestration platforms, RAG retrieval systems with LLM evaluation, and production observability tooling (Datadog + Vertex AI). Proven impact: 80% accuracy improvement in regulated ML systems; 70% faster workflows via production agentic AI. Comfortable leading applied AI investments end-to-end from architecture through deployment and monitoring.

CORE COMPETENCIES

Applied AI in Regulated Environments

Agentic AI Architecture (LangChain, LangGraph)

RAG Pipelines & Vector DB Design

LLMOps & Production Observability

HIPAA-adjacent Compliance Design

Multi-Agent Orchestration (Google ADK)

GCP / Vertex AI / Cloud Run

MLflow & Datadog Monitoring

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven data validation and automation systems for BMW service centres — achieving **80% accuracy improvement** in a regulated, high-volume production environment with strict GDPR and data integrity requirements.
- Operated across healthcare, mobility, and industrial domains — building compliance-first ML pipelines with normalisation rules, label quality checks, and automated anomaly detection; directly relevant to HIPAA-adjacent AI platform design.
- Implemented observability via structured error handling and logging, reducing recurring production failures; demonstrated reliability-first engineering mindset applicable to clinical AI systems.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

Multi-Agent Productivity Assistant — Agentic AI Platform Gen AI Academy APAC · 2026

Architected multi-agent AI system from scratch: Orchestrator Agent routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, real-time tool-use via MCP servers, parallel execution, and AlloyDB AI persistent memory. Deployed on Cloud Run. Demonstrates the agentic AI platform architecture Rula requires.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

TravelAI — RAG Pipeline with LLM Evaluation

Kaggle Capstone · 2025

End-to-end RAG implementation: ingestion pipelines, chunking and embedding policies, hybrid vector search, LLM grounding, and controlled retrieval relevance experiments. Demonstrates RAG-at-scale engineering — the retrieval and grounding patterns central to clinical AI assistants.

RAG · Pinecone-style vector DB · Embeddings · Semantic search · LLM grounding · Python

AI SRE Copilot — LLMops Observability

DEVPOST Hackathon · 2025

LLMops observability infrastructure integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis of production AI systems. FastAPI on Cloud Run. Demonstrates real-time LLM monitoring — cost tracking, safety monitoring of AI outputs.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMops

WorkflowEnv — OpenEnv RL Environment for LLM Agents

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

OpenEnv-compliant RL environment for LLM workflow automation agents — 7-criterion partial-credit grader, telemetry logging, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.**

FastAPI · Docker · HF Spaces · OpenEnv · LLM agent evaluation

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany

2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University

2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMops, OpenAI API, LangChain, Pinecone, Embeddings, Production AI)

2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive

2024–2025

SKILLS

Agentic AI: LangChain, LangGraph, Google ADK, multi-agent orchestration, tool-use, stateful agent design, ReAct pattern

LLMs & RAG: Gemini 2.5, OpenAI API, RAG pipelines, vector DBs (Pinecone), embeddings, prompt engineering, LLM evaluation

MLOps: MLflow, Datadog, GCP Vertex AI, Cloud Run, Cloud Build CI/CD, Docker **Languages:** Python, FastAPI, Node.js, SQL, Bash