

Darshan Linge Gowda

darshan1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI/ML Engineer specialising in RAG pipelines, document AI, and end-to-end LLM system delivery — directly applicable to building AI-powered technical report generation for insurance claims. Built production RAG systems with controlled retrieval experiments, multi-agent AI workflows deployed on GCP Cloud Run with real users, and an OpenEnv RL evaluation environment validated at 0.96/1.0. Three years of enterprise data engineering across BMW, healthcare, and industrial domains — normalisation rules, label quality checks, large-scale schema migrations — adds the data handling rigour that production document AI requires. Strong command of Python, LangChain, LangGraph, FastAPI, vector DBs, and Gemini/OpenAI APIs.

CORE COMPETENCIES

RAG Pipelines & Document AI

LLM API Integration (Gemini, OpenAI)

End-to-End LLM System Delivery

Vector DBs & Embedding Pipelines

LLM Evaluation & Retrieval Experiments

Unstructured Data & ETL Pipelines

FastAPI & Docker / Cloud Run

Python & LangChain / LangGraph

WORK EXPERIENCE

PICA GmbH2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven information extraction and data validation systems for BMW service centres — applying feature engineering, normalisation rules, label quality checks, and anomaly detection on structured and semi-structured datasets, delivering **80% accuracy improvement** in production.
- Designed compliance-first ML pipelines across healthcare, mobility, and industrial domains in regulated, high-volume environments with strict data governance — directly applicable to document-centric AI platform engineering.
- Implemented structured error handling and logging to reduce recurring production failures through systematic root-cause analysis.

AAM IT GmbH / Vorwerk SwitzerlandJan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk’s Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

AdremesMay–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows. Developed backend validation logic and session-handling for a cloud-based ad-analytics platform.

SELECTED AI/ML PROJECTS

TravelAI — RAG Pipeline with Retrieval EvaluationKaggle Capstone · 2025

End-to-end RAG system: heterogeneous data ingestion, chunking and embedding pipeline, hybrid vector search, LLM grounding, and controlled retrieval experiments measuring relevance and consistency across chunking strategies, embedding models, and retrieval configurations. Directly demonstrates the RAG engineering and LLM evaluation skills required to build production-quality technical report generation systems.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

AI Invoice Agent — Prototype to Production LLM System

Gen AI Academy APAC · 2026

Took a conversational multi-agent AI from concept to production: Gemini 2.5 Flash powering invoice creation, payment tracking, and analytics. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) in a live system. **Impact: ~70% faster workflows.** Full-stack ownership: React + Vite, Node.js, Docker, Cloud Run.

Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

WorkflowEnv — LLM Agent Evaluation Framework

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment for LLM workflow automation agents — 7-criterion partial-credit grader, telemetry logging, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.** Demonstrates systematic LLM evaluation and production deployment capability.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · telemetry logging

Multi-Agent Productivity Assistant — Production Agentic System

Gen AI Academy APAC · 2026

Architected multi-agent AI from scratch: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, real-time tool-use via MCP servers, parallel execution, AlloyDB AI persistent memory. Deployed on Cloud Run with real users.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Semantic Search)2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

LLMs & RAG: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM grounding, retrieval evaluation

Data Engineering: ETL pipelines, XML/JSON schema transformation, anomaly detection, data validation, unstructured data processing, Python, SQL

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, Azure/AWS familiarity **Languages:** Python, FastAPI, Node.js, SQL, Bash, C#