

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Engineer with a track record of building and shipping production-grade agentic AI pipelines, LLM evaluation frameworks, and RAG systems — exactly the scope of StackHawk’s AI Software Engineer role. Built an OpenEnv-compliant RL evaluation environment validated at 0.96/1.0 (Meta × Hugging Face 2026), two agentic AI systems deployed on GCP Cloud Run with real users, and a multimodal AI observability copilot. Strong command of the full stack from prototype to production: LangChain, LangGraph, FastAPI, Docker, Cloud Run, Gemini/OpenAI APIs, and LLM evaluation design. Three years of enterprise software engineering in regulated environments adds the production reliability discipline that security-critical AI features demand.

CORE COMPETENCIES

Agentic AI Pipeline Design

LLM Evaluation Framework Design

RAG & Semantic Retrieval

Production AI Features (Cloud Run)

LangChain & LangGraph

FastAPI & Docker CI-CD

LLM API Integration (Gemini, OpenAI)

Multi-Agent Orchestration

WORK EXPERIENCE

PICA GmbH

2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML systems for BMW service centres — from data pipelines and model design through production deployment and monitoring — delivering **80% accuracy improvement** in automated invoice generation.
- Built production backend systems across healthcare, mobility, and industrial domains in regulated, high-volume environments with strict data governance requirements.
- Implemented structured observability and automated failure detection; reduced recurring production failures through systematic root-cause analysis — directly applicable to reliable AI feature engineering in security-critical products.

AAM IT GmbH / Vorwerk Switzerland

Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk’s Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes

May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — LLM Agent Evaluation Framework

OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL environment for evaluating LLM workflow automation agents — FastAPI server, 7-criterion partial-credit grader, strict telemetry logging, Docker deployment to HF Spaces. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.** Demonstrates systematic LLM evaluation design at production quality — directly applicable to StackHawk’s AI feature evaluation work.

FastAPI · Docker · HF Spaces · OpenEnv · LLM agent evaluation · telemetry logging

Multi-Agent Productivity Assistant — Agentic AI Pipeline

Gen AI Academy APAC · 2026

Built and shipped a multi-agent agentic pipeline from scratch: Orchestrator Agent routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, real-time tool-use via MCP servers, parallel agent execution, AlloyDB AI persistent memory. Deployed on Cloud Run with real users.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Production LLM Reliability Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI for invoice creation and analytics; diagnosed and resolved a critical production LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Full-stack: React + Vite, Node.js, Docker, Cloud Run.
Gemini 2.5 Flash · LangChain · React + Vite · Node.js · Docker · Cloud Run

AI SRE Copilot — Multimodal LLM Application DEVPOST Hackathon · 2025

Multimodal AI application integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis of production incidents. FastAPI on Cloud Run.
Demonstrates LLM pipeline design for complex, data-rich operational inputs.
Vertex AI · Gemini · Datadog · FastAPI · Cloud Run

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021
Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

LLM & Agents: Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, multi-agent orchestration, prompt engineering, LLM evaluation, fine-tuning

AI Systems: Agentic pipelines, RAG, vector DBs (Pinecone), embeddings, semantic search, ReAct pattern, RL environment design

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, MLflow, Datadog

Languages: Python, FastAPI, Node.js, TypeScript, SQL, Bash