

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

AI Software Engineer with hands-on experience building agentic pipelines, automated evaluation systems, and production multi-agent backends — the engineering foundation StackHawk’s AI-powered security testing requires. Shipped an RL evaluation environment with automated agent scoring and Docker deployment (WorkflowEnv, 0.96/1.0), production multi-agent systems with real-time MCP tool-use and parallel execution on Cloud Run, and a FastAPI RAG pipeline with systematic retrieval experiments. Strong command of Python, FastAPI, LangChain, LangGraph, and GCP Cloud Run — backed by three years of enterprise backend engineering in production-critical environments where automated failure detection and reliability under high-volume workloads are non-negotiable.

CORE COMPETENCIES

Agentic Pipeline Design & Deployment

Automated Evaluation & Scoring Systems

Multi-Agent Orchestration (Google ADK)

LangChain & LangGraph

FastAPI & Python Backend Engineering

Docker & GCP Cloud Run CI/CD

RAG Pipelines & Retrieval Systems

Production LLM Reliability

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven automation systems for BMW service centres — end-to-end pipeline ownership from model design through production deployment, delivering **80% accuracy improvement** in automated failure detection.
- Engineered structured observability and automated testing/detection pipelines across healthcare, mobility, and industrial domains — the production reliability discipline that AI-powered security testing pipelines demand.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk’s Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

WorkflowEnv — Automated Agent Evaluation Pipeline OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an automated agent evaluation pipeline — exactly the kind of scoring and testing infrastructure StackHawk’s AI security agents require. 7-criterion partial-credit grader scoring agent task completion, instruction adherence, constraint satisfaction, and output quality. Full telemetry logging, automated scoring pipeline connecting agent outputs to reward signals, Docker deployment to HF Spaces. Both Phase 1 and Phase 2 validated. **Avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL evaluation design · automated agent scoring · telemetry · reward pipeline

Multi-Agent Productivity Assistant — Agentic System with Real-Time Tool Calling Gen AI Academy APAC · 2026

Architected a production agentic system with real-time MCP tool integrations: Orchestrator Agent routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel agent execution, Google Calendar MCP + Gmail MCP live tool-use, AlloyDB AI persistent session memory. Deployed on Cloud Run with CI/CD and real users. Demonstrates production-grade agentic pipeline architecture at the level StackHawk’s AI engine requires.

Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

AI Invoice Agent — Full-Stack Production LLM Application Gen AI Academy APAC · 2026

Shipped a conversational multi-agent AI from prototype to production — FastAPI + Node.js backend, Gemini 2.5 Flash LLM orchestration via LangChain, React + Vite frontend. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows.** Docker + Cloud Run deployment end-to-end.

Gemini 2.5 Flash · LangChain · React + Vite · FastAPI · Node.js · Docker · Cloud Run

TravelAI — RAG Pipeline with Systematic Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG system with controlled experiments measuring retrieval accuracy across chunking strategies, embedding models, and retrieval configurations — the systematic evaluation methodology applicable to testing AI response quality in security tool workflows.

RAG · Vector DB · Pinecone · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

EDUCATION

- M.Sc. Information & Communication Systems** · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering** · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

- DataCamp** — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

- Agentic AI:** LangChain, LangGraph, Google ADK, multi-agent orchestration, ReAct pattern, MCP tool-use, parallel execution, automated agent evaluation, reward pipeline design, telemetry

LLMs & RAG: Gemini 2.5/2.0, OpenAI API, RAG pipelines, vector DBs (Pinecone), embeddings, semantic search, LLM evaluation, prompt engineering

Cloud / Infra: GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, AlloyDB AI, Datadog

Languages: Python, FastAPI, Node.js, TypeScript, React, SQL, Bash