

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

PROFESSIONAL SUMMARY

Full-Stack Software Engineer with hands-on experience shipping production AI systems from prototype to real users — exactly what Supermove's growing AI feature set requires. Shipped a conversational multi-agent AI with **~70% workflow acceleration** in production, a production multi-agent orchestration system on Cloud Run with real users, and an LLM evaluation environment validated at 0.96/1.0. Strong command of Python, React, FastAPI, Node.js, SQL, and LangChain — the full-stack AI engineering skills to build and own features end-to-end. Three years of enterprise backend engineering in high-volume regulated environments adds the production discipline that a fast-moving startup demands.

CORE COMPETENCIES

Full-Stack Engineering (Python + React)

Production AI Features in SaaS Products

Conversational AI & LLM Orchestration

LangChain & LangGraph

FastAPI & Python Backend

React & TypeScript Frontend

Multi-Agent Orchestration

Docker & GCP Cloud Run

WORK EXPERIENCE

PICA GmbH

2018 – 2021

Software Developer · Munich, Germany

- Owned end-to-end ML backend systems for BMW service centres — from data pipelines and feature engineering through production deployment and monitoring, delivering **80% accuracy improvement**.
- Built production backend systems across healthcare, mobility, and industrial domains; maintained reliability under enterprise-scale workloads — the full-stack production discipline a fast-moving startup requires.

AAM IT GmbH / Vorwerk Switzerland

Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes

May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

SELECTED AI/ML PROJECTS

AI Invoice Agent — Full-Stack Production AI SaaS Feature

Gen AI Academy APAC · 2026

Shipped a full-stack conversational multi-agent AI from prototype to production — the exact profile of an AI feature in a SaaS product. React + Vite frontend, FastAPI + Node.js backend, Gemini 2.5 Flash LLM orchestration via LangChain: intent routing, invoice creation, payment tracking, and analytics in a single conversational UI. Diagnosed and resolved a critical LLM reliability issue (API version mismatch) mid-flight. **Impact: ~70% faster workflows for real users.**
Gemini 2.5 Flash · LangChain · React + Vite · FastAPI · Node.js · Docker · Cloud Run

Multi-Agent Productivity Assistant — Production Agentic System

Gen AI Academy APAC · 2026

Architected and shipped a multi-agent AI system from scratch: Orchestrator routing intent to four specialist sub-agents (Calendar, Task, Notes, Email) via Gemini 2.0 Flash + Google ADK, parallel execution resolving complex requests in a single turn, real-time MCP tool-use, AlloyDB AI persistent session memory. Cloud Build CI/CD, Cloud Run. Real users in production.
Gemini 2.0 Flash · Google ADK · LangGraph · MCP · FastAPI · Cloud Build CI/CD · Cloud Run · AlloyDB AI

WorkflowEnv — LLM Evaluation Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL evaluation environment — 7-criterion partial-credit grader, automated scoring pipeline, telemetry logging, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0.** Demonstrates the systematic AI evaluation and production reliability thinking Supermove's AI features require.

FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · automated scoring · telemetry

TravelAI — RAG Pipeline & Data Integration Kaggle Capstone · 2025

End-to-end RAG system: heterogeneous data ingestion, chunking and embedding, hybrid vector search, LLM grounding, and controlled retrieval experiments measuring response quality across configurations.

RAG · Vector DB · Pinecone · Embeddings · Semantic search · LLM grounding · Python

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMOps, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

Full-Stack: Python, FastAPI, React, TypeScript, Node.js, SQL, REST APIs, GraphQL

AI & LLMs: LangChain, LangGraph, Google ADK, Gemini 2.5/2.0, OpenAI API, RAG pipelines, vector DBs (Pinecone), prompt engineering, multi-agent orchestration

Cloud / Infra: GCP, Cloud Run, Cloud Build CI/CD, Docker, AlloyDB AI, HF Spaces, Vertex AI

Observability: Datadog, MLflow, structured logging, LLMOps, automated failure detection