

Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda |

Bangalore, India · Open to Remote (Worldwide)

PROFESSIONAL SUMMARY

ML/LLMOps Engineer with hands-on experience building CI/CD pipelines for LLM deployment, containerised model infrastructure on GCP Cloud Run, and production observability systems using Datadog and Vertex AI. Delivered reusable CI/CD workflows for model training and deployment across multiple production systems — including a multi-agent orchestration platform and an LLM observability copilot. Proven impact: 80% accuracy improvement in production ML systems for automotive clients; infrastructure reliability-first approach from 3 years in regulated German enterprise environments. Experienced with MLflow for experiment tracking, vector databases for semantic retrieval, and Docker-based deployment at scale.

CORE COMPETENCIES

CI/CD for LLM Training & Deployment

Vector Database Management

Model Observability & Monitoring

MLflow & Model Registry

Docker & Container Orchestration

GCP Cloud Run & Cloud Build

LLMOps Infrastructure Design

Production LLM Reliability Engineering

WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Delivered **80% accuracy improvement** and 50% processing time reduction in ML-driven SAP invoice automation for BMW service centres — full pipeline ownership from data validation through production deployment and monitoring.
- Built automated CI/CD workflows and validation pipelines for ML models operating at high volume across regulated enterprise environments (healthcare, automotive, industrial).
- Implemented observability via structured error handling, logging, and automated failure detection — reducing recurring production failures; directly applicable to LLMOps monitoring.
- Operated across multi-tenant environments with strict data governance and compliance requirements.

AAM IT GmbH 2018

Software Developer (Contract) · Munich, Germany

- Designed and executed large-scale data migration with version-controlled schema evolution; built automated ETL pipelines via Jenkins and Bash — demonstrates model versioning and pipeline automation skills.
- Developed C# APIs for multi-interface integrations; provided cross-timezone production support ensuring pipeline stability.

SELECTED AI/ML PROJECTS

WorkflowEnv — OpenEnv RL Environment for LLM Agents OpenEnv AI Hackathon · Meta x Hugging Face · 2026

Built and deployed an OpenEnv-compliant RL environment for LLM workflow automation agents. Implemented FastAPI server, 7-criterion partial-credit grader, Docker deployment to HF Spaces, and strict telemetry logging. Phase 1 and Phase 2 both validated — **avg score 0.96/1.0, 3/3 tasks passing.**

FastAPI · Docker · HF Spaces · OpenEnv · RL environment design · LLM agent evaluation · telemetry logging

AI SRE Copilot — LLMops Observability Infrastructure DEVPOST Hackathon · 2025

Built LLMops observability infrastructure integrating Datadog metrics/logs with Vertex AI (Gemini) for automated root-cause analysis. FastAPI service deployed serverlessly on Cloud Run with structured AI insights and real-time performance monitoring — directly models the observability tooling TRM Labs needs.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMops observability · Python

Multi-Agent Productivity Assistant — CI/CD for Agentic Systems Gen AI Academy APAC · 2026

Designed and deployed a multi-agent system with reusable CI/CD pipeline (Cloud Build + Docker) for containerised deployment on Cloud Run. AlloyDB AI for persistent memory. Demonstrates automated model deployment and versioning infrastructure at scale.

Gemini 2.0 Flash · Google ADK · Cloud Build CI/CD · Docker · Cloud Run · AlloyDB AI

AI Invoice Agent — Production LLM Reliability Engineering Gen AI Academy APAC · 2026

Diagnosed and resolved critical production LLM reliability issue (model/API version mismatch) in a live multi-agent system — achieving stable low-latency responses at scale. **Impact: ~70% faster workflows.** Demonstrates the production reliability engineering mindset central to LLMops.

Gemini 2.5 Flash · LangChain · Docker · Cloud Run · CI/CD · LLM reliability

Generative AI Search Agent — Vector DB & Semantic Retrieval Personal Project · Ongoing

LLMops-focused semantic search agent: vector database management, embedding pipeline, prompt chaining, contextual memory. Containerised with Docker/CI-CD. **~40% faster resolution vs. keyword search.**

LangChain · Vector DB · Embeddings · OpenAI API · Docker · CI/CD · MLflow

EDUCATION

M.Sc. Information & Communication Systems · TU Chemnitz, Germany 2017 – 2021

Thesis: IEEE-published — ML pipeline for large-scale geospatial data processing adopted by a Galileo Map Service Provider for production use.

B.E. Electronics & Communication Engineering · Visvesvaraya Technological University 2013 – 2017

CERTIFICATIONS

DataCamp — Associate AI Engineer for Developers (LLMops, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems) 2025

MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive 2024–2025

SKILLS

LLMops / MLOps: CI/CD for model deployment, MLflow, Langfuse-adjacent (Datadog), vector DBs, model registries, Docker, Kubernetes-ready containerisation

Cloud Infrastructure: GCP, Vertex AI, Cloud Run, Cloud Build, serverless deployment, multi-region scaling

LLMs & GenAI: Gemini 2.5, OpenAI API, LangChain, LangGraph, RAG, prompt engineering, multi-agent orchestration

Languages: Python, FastAPI, Node.js, SQL, Bash