

# Darshan Linge Gowda

darshanl1711@gmail.com | linkedin.com/in/darshan-lingegowda-ai | github.com/DarshanLingegowda | Bangalore, India · Open to Remote (USA)

## PROFESSIONAL SUMMARY

AI Engineer specialising in LLM evaluation, RAG pipelines, and production-grade model deployment — the core of Valence's Applied AI stack. Built an OpenEnv-compliant LLM evaluation framework validated at 0.96/1.0 (Meta × Hugging Face 2026), a RAG concierge with controlled retrieval experiments across chunking and embedding strategies, and two production multi-agent systems deployed on GCP Cloud Run. Strong command of LLM component development and evaluation, data integration pipelines, cloud deployment with observability (MLflow, Datadog), and rapid prototyping from concept to production. Three years of enterprise software engineering adds the reliability and cross-functional collaboration discipline that Valence's coaching platform requires.

## CORE COMPETENCIES

LLM Evaluation & Component Development

RAG Pipelines & Data Integration

Production Model Deployment & Observability

LangChain & LangGraph

GCP / Vertex AI / Cloud Run

Multi-Agent Orchestration (Google ADK)

MLflow & Datadog Monitoring

Rapid Prototyping to Production

## WORK EXPERIENCE

PICA GmbH 2018 – 2021

Software Developer · Munich, Germany

- Built ML-driven data validation and automation systems for BMW service centres — designing pipelines from feature engineering and model evaluation through deployment and monitoring, delivering **80% accuracy improvement**.
- Implemented structured observability via logging and automated failure detection; reduced recurring production failures through systematic root-cause analysis — the reliability mindset Valence's continuous model improvement requires.
- Collaborated cross-functionally with domain experts to translate business requirements into production ML solutions across healthcare, mobility, and industrial domains.

AAM IT GmbH / Vorwerk Switzerland Jan–Oct 2018

Software Developer (Contract) · Munich, Germany

- Engineered C# backend APIs, SOAP-based integrations, and SQL Server data migration pipelines for Vorwerk's Thermomix smart appliance platform; led large-scale relational-to-XML schema transformation ensuring high data integrity across distributed systems.
- Built Python automation for XML dataset validation, anomaly detection, and ETL optimisation using Jenkins, Bash, and PHP — significantly reducing manual intervention and improving pipeline reliability.

Adremes May–Oct 2017

Software Developer (Trainee) · Hamburg, Germany

- Reduced manual ad scheduling workload by **40%** building automated backend tooling for radio ad slot booking — C# MVC application integrated into the media planning platform, replacing spreadsheet-driven workflows.

## SELECTED AI/ML PROJECTS

WorkflowEnv — LLM Evaluation Framework OpenEnv AI Hackathon · Meta × Hugging Face · 2026

Designed and deployed an OpenEnv-compliant RL evaluation environment for LLM workflow agents — 7-criterion partial-credit grader with strict telemetry logging, automated scoring pipeline, Docker deployment to HF Spaces. Both phases validated. **Avg score 0.96/1.0, 3/3 tasks passing.** Demonstrates systematic LLM evaluation, observability tooling, and production deployment — the continuous improvement infrastructure Valence's AI stack requires.  
FastAPI · Docker · HF Spaces · OpenEnv · LLM evaluation · telemetry logging · automated scoring

TravelAI — RAG Pipeline with Retrieval Evaluation Kaggle Capstone · 2025

End-to-end RAG system: heterogeneous data ingestion, chunking and embedding pipeline, hybrid vector search, LLM grounding, and controlled retrieval experiments measuring relevance and consistency across embedding models and retrieval configurations. Demonstrates the data integration and LLM component evaluation skills central to Valence's contextual AI coaching.

RAG · Vector DB · Embeddings · Semantic search · LLM grounding · Retrieval evaluation · Python

AI SRE Copilot — LLMops Observability DEVPOST Hackathon · 2025

LLMops observability infrastructure integrating Datadog with Vertex AI (Gemini) for automated root-cause analysis. FastAPI on Cloud Run. Demonstrates real-time LLM monitoring and structured AI insights — the model deployment observability Valence's reliability requirements demand.

Vertex AI · Gemini · Datadog · FastAPI · Cloud Run · LLMops

Multi-Agent Productivity Assistant — Production Agentic System Gen AI Academy APAC · 2026

Architected multi-agent AI from scratch: Orchestrator routing intent to four specialist sub-agents via Gemini 2.0 Flash + Google ADK, parallel execution, AlloyDB AI persistent memory. Deployed on Cloud Run with real users. Demonstrates full-stack AI prototyping to production.

Gemini 2.0 Flash · Google ADK · LangGraph · Cloud Build CI/CD · Cloud Run · AlloyDB AI

EDUCATION

- M.Sc. Information & Communication Systems · TU Chemnitz, Germany2017 – 2021
- Thesis: IEEE-published — large-scale geospatial ML pipeline adopted by a Galileo Map Service Provider for production use.
- B.E. Electronics & Communication Engineering · Visvesvaraya Technological University2013 – 2017

CERTIFICATIONS

- DataCamp — Associate AI Engineer for Developers (LLMops, OpenAI API, LangChain, Pinecone, Embeddings, Production AI Systems)2025
- MIT-IDSS — Data Science & Machine Learning · Kaggle Generative AI & AI Agents Intensive2024–2025

SKILLS

**LLMs & Eval:** Gemini 2.5/2.0, OpenAI API, LangChain, LangGraph, Google ADK, LLM evaluation, prompt engineering, fine-tuning, RLHF concepts, RAG pipelines

**MLOps:** MLflow, Datadog, structured observability, vector DBs (Pinecone), embeddings, semantic search, automated scoring pipelines

**Cloud / Infra:** GCP, Vertex AI, Cloud Run, Cloud Build CI/CD, Docker, FastAPI, HF Spaces, AWS/Azure familiarity

**Languages:** Python, FastAPI, Node.js, SQL, Bash